Vega training from correctness feedback

Gray: fixed computation or data. Green: learned adjustment and its loss.

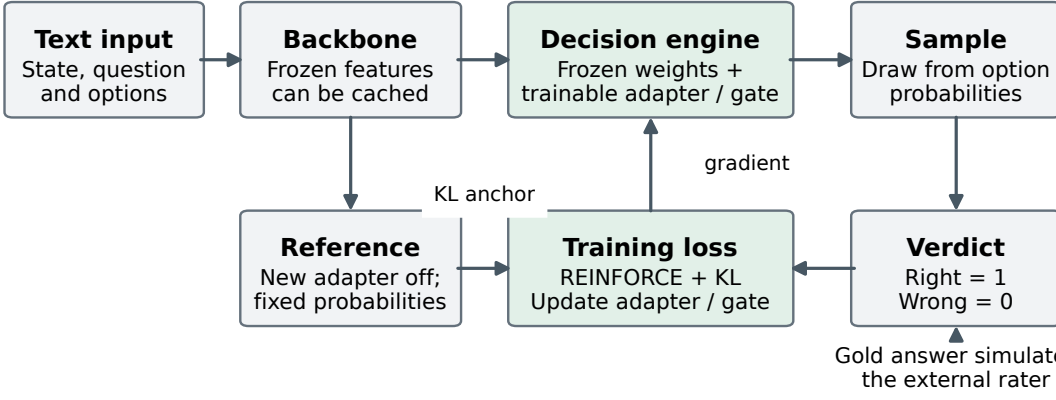

One pass	One new sampled answer and verdict per training question.
Eight passes	Repeat each training question with fresh samples and verdicts.

After training: held-out verdicts → temperature fitting → optional model-level switch.