Peekaboo: A Study of Right / Wrong Feedback for Small Text Decision Models

Abhijith Neil Abraham
VitalOps

November 2026

Abstract

We study how much of supervised adaptation small decision models recover from right / wrong feedback on their own answers. The evaluation covers text decision tasks on two released model families, Vega and Laya, using benchmark gold answers to simulate feedback. The primary setting collects one verdict per training row; an eight-pass setting measures what repeated feedback adds. On DecisionBench at 0.8B, one on-policy pass improves Vega by +6.7 accuracy points, while reusing one logged sampled verdict per row for eight optimisation passes gives +8.3. Eight passes with fresh verdicts give +11.3, against +13.2 for a matched labelled adapter with early stopping. The logged arm remains significantly below repeated feedback at 0.8B, although it also changes the collection policy and training objective. Repeated feedback often identifies the correct answer, and the expected REINFORCE reward gradient weights the supervised gradient by the model's current probability of that answer. Label-prior controls recover much of the gain on typed-decisions and the public task menu, but not on DecisionBench and emotion. On Laya, training the encoder with verdicts improves the released model, while matched labelled training performs better on typed-decisions and emotion; those encoder comparisons have one seed on each of these tasks. Temperature scaling accounts for much of the calibration improvement, and benchmark shortcuts and training-pool overlap limit generalisation. The contribution is a controlled empirical comparison of feedback budgets and matched adaptation methods, not a new estimator or evidence of performance with real users or multimodal inputs. Code and per-row records are released.

1 Introduction

Vega [37] and Laya [6] are small models built for one job: given a situation and a typed question (pick one of these options, yes or no, a rating), return a probability for each option in one forward pass, fast enough and calibrated enough to sit inside an agent loop and route the uncertain cases to a person. Both ship with a recipe for getting better at a user's task, and both recipes want labelled rows. Vega's is a per-task classifier fitted to its frozen features (test-time training, TTT) or an adapter trained the same way; Laya's is a whole-model fine-tune, and its released `laya-typed-decisions` checkpoint is that recipe applied to its headline benchmark.

A deployed model may receive correctness feedback more readily than complete labels: the suggested tool call was cancelled, the flagged message was fine after all, the escalated ticket did need escalating. Such events can motivate a feedback signal, but their interpretation as correctness is application-dependent; our experiments instead simulate unambiguous verdicts from gold answers. Learning a classifier from that signal is the bandit multiclass problem [7, 9, 26], and improving a deployed NLP system from binary feedback simulated from gold labels is the protocol of Gao, Choi and Artzi [14]. We use this established setting to compare adaptation methods and feedback budgets.

We ask: for these two models, on the benchmarks their own authors use, how much of the labelled-data recipe does right / wrong feedback recover, under what protocol, and what are the gains made of? The protocol matters: eight on-policy verdicts identify the gold answer for most rows, and the expected REINFORCE reward gradient is the cross-entropy gradient reweighted by the model's own probability of the gold answer [10, 59]. We therefore distinguish one rating per row from repeated feedback, compare matched training recipes, and measure how much of the gain comes from label priors or benchmark artifacts.

Scope and research questions. We evaluate text and structured states processed through the text input path. Vega supports image inputs, but this study neither trains on images nor evaluates multimodal transfer or retention. We ask (i) how much adaptation one verdict per training row supports, (ii) what repeated feedback adds relative to matched labelled training, and (iii) how much of the gain remains after accounting for label priors, benchmark shortcuts and confidence calibration. These are empirical questions about the evaluated models, not claims of a new reinforcement-learning algorithm.

The method, Peekaboo, has three parts: an add-on next to the untouched released model (a rank-8 gated LoRA adapter on Vega's decision engine; a copy of Laya's decision head, or of its head and encoder), trained by REINFORCE on the verdict of a sampled answer with a running baseline and a KL term towards the released probabilities; a single temperature refitted from verdicts alone; and a model-level switch that keeps the add-on when its held-out best-answer accuracy exceeds the released model's. The switch is an estimate from a finite calibration set, not a guarantee against regressions. Our evaluation includes:

- the clean feedback arms are run with one pass in which each training row is rated exactly once and with eight passes that collect fresh verdicts, and so is the answer-key adapter on Vega, at matched epochs; on Laya the head is run for one pass and the encoder for eight and one;
- simulated feedback under deployment-inspired collection policies: the released model's own best or sampled answers, logged and rated once and learned off-policy, and (under eight passes) ratings that are missing, flipped, or skewed towards complaints;
- matched baselines and controls: the same add-on trained on the labels, with and without early stopping; a label-prior baseline with no input features, trained from the same verdicts; the released model with the same feedback-fitted temperature; TTT under both of VegaML's selection rules; for Laya, the matched in-domain answer-key fine-tune with the encoder unfrozen;
- intervals that resample training seeds as well as rows, clustered by task or case, with Holm adjustment within each family of comparisons and a non-inferiority test at a stated margin; proper scores next to ECE;
- disclosure of what the benchmarks contain: a TruthfulQA task whose gold answer is always first, a DecisionBench task whose gold is readable from the option structure, and a shipped Vega adapter that was fitted with labels on the very rows that serve as its typed-decisions feedback pool.

What we found. With one verdict per training row, one on-policy pass improves Vega on DecisionBench by +6.7 points at 0.8B and +3.2 at 4B. Reusing a single logged sampled verdict per row for eight optimisation passes gives +8.3 and +4.3. Eight passes that collect fresh verdicts give +11.3 and +5.7; the logged arm remains 2.9 points behind at 0.8B, a difference that survives Holm adjustment. This comparison is consistent with a benefit from additional feedback, but does not isolate information from the change in collection policy and off-policy objective. Repeated verdicts are particularly informative: on DecisionBench each row produces 4.88 correct-answer events on average at 0.8B, each identifying its label. Eight-pass feedback therefore approaches the labelled adapter's point estimates on most Vega benchmarks without establishing equivalence. A per-task label prior recovers much of the gain on typed-decisions and the public menu; DecisionBench and emotion provide clearer evidence of improvement beyond prior correction. On Laya, encoder training with verdicts improves the released model, but matched labelled training wins on typed-decisions and emotion by 6.5 and 7.9 points. These two encoder comparisons have one seed each. The calibration,

KL and switch controls identify useful components and failure modes, rather than guarantees of deployment safety. All one-rating descriptions concern training rows; calibration and model selection use additional held-out feedback.

2 Related work

Learning a classifier from right / wrong. Multiclass learning in which the learner is told only whether its prediction was correct goes back to the Banditron [26]; Crammer and Gentile [7] evaluated it on text classification, and Daniely and Halbertal [9] bound the price of bandit information at $O(k \log k)$ more mistakes in the realisable case. Wang and Qiao [52] train a neural classifier online from correctness feedback with an importance-weighted cross-entropy and derive conformal sets from the same feedback. Learning a structured predictor from the reward of its own output is bandit structured prediction [28, 38, 48], shown to work with real human ratings [30] and, at eBay, with implicit feedback where explicit ratings were too noisy [29]; the off-policy branch learns from logged feedback with importance weights [12, 24, 33, 51].

Improving a deployed system from its users. The protocol closest to ours is Gao, Choi and Artzi [14]: binary rewards simulated from gold labels, on-policy policy-gradient learning from the model's own answers, supervised training as the reference line, and the finding that the initial model's quality decides whether learning works. Its follow-up collected real feedback over nine deployment rounds [15]. Campos et al. [3] and Hancock et al. [20] improve deployed QA and dialogue systems from binary signals. Our contribution relative to this line is the object of study, two released typed-decision models with calibrated probabilities, the comparison against their vendors' own labelled-data recipes, and the matched one-pass and logged-feedback arms.

What a 0/1 reward does to the update. With a single gold answer y and reward $r = \mathbf{1}[a = y]$, the expected REINFORCE update is $\mathbb{E}_{a\sim\pi}[(r-b)\nabla \log \pi(a)] = \pi(y)\,\nabla \log \pi(y) = \nabla \pi(y)$: cross-entropy on the gold answer, reweighted by the policy's own probability of it. This is the mean-absolute-error objective of the label-noise literature [16, 59], known to under-fit rows the model already gets badly wrong, and the identity Davis and Recht [10] give for binary-reward RL in general. Filtered self-training [11, 55, 57] is the same mechanism without the negative term: every correct sample is a supervised step on the revealed label. Yue et al. [56] show that RL with verifiable rewards mostly re-ranks what the base model could already produce; Choshen et al. [4] and Kiegeland and Kreutzer [27] make the same point for bandit NMT and attribute much of the gain to the sharpening of the output distribution. Our fix / break analysis (Section 6.3) is a direct check of these predictions on our models.

KL-regularised fine-tuning and forgetting. A KL penalty to a reference policy is standard in RLHF [5, 39, 50, 63] and offline RL for dialogue [23]; DPO [43] and KTO [13] fold it into a supervised objective, KTO from binary desirable / undesirable labels, and REINFORCE with a baseline remains competitive with more elaborate estimators [1]. Shenfeld et al. [47] show that on-policy RL forgets less than supervised fine-tuning because it stays closer to the base distribution, which bears on how Laya's whole-model fine-tune and our encoder-level add-on behave away from their training benchmark.

Few-shot heads against a zero-shot model. A linear probe fitted on a few labelled rows can be worse than the zero-shot model it replaces; CLIP needed about four shots per class to match its own zero-shot accuracy [42], and anchoring the fine-tuned model to the zero-shot one is the standard remedy [32, 54]. VegaML's TTT head is such a probe, and the rule that decides when it replaces the engine changed in vegaml 0.12.3 to require beating the engine as well as chance; we report both rules.

Label priors and label shift. Part of what any adapter learns from a benchmark's rows is its label distribution and answer-position bias. Contextual calibration [60], batch calibration [62], prior adjustment [34, 45] and the choices-only baselines of multiple-choice evaluation [2, 61] all measure this; our label-prior baseline (Section 6.3) is the feedback-trained version.

Calibration after policy gradient. The expected 0/1 reward of a sampled answer is the linear score, which is improper [17]: it is maximised by putting all the mass on the argmax, and under a KL term of weight β the optimum raises the log-odds of a reliably right answer by up to $1/\beta$ nats. Over-confidence after RL fine-tuning is documented for RLHF policies and fixed with a single temperature [25], and binary-correctness RL hurts calibration unless a proper score is added to the reward [8]; Laya's own RLCD uses proper scoring rules for this reason. Fitting a temperature from correctness alone is top-label calibration [19]: ECE itself [18, 36] consumes nothing but (confidence, right / wrong) pairs. Our "honesty fix" is that, written for sampled rather than argmax answers; we no longer present it as calibration without labels in any new sense. ECE is a biased, non-proper estimate [31, 44], so we report NLL, Brier and AURC beside it.

3 The two decision models, and what they already know

Vega. Vega [37] reads the situation, question and options with a frozen Qwen3.5 backbone [40, 41] of 0.8B or 4B parameters, pools hidden states from two intermediate layers, and hands them to a small decision engine (14.3M parameters at 0.8B, 30.9M at 4B) that the authors call a physics engine: each option is an energy well in a 64-dimensional landscape, a particle is released, and $P(a) = \mathrm{softmax}(-E_a/\tau(s))$ with a per-question temperature computed from the particle's trajectory. The backbone is not updated during Vega's task adaptation; the engine and two gated adapters were, on an undisclosed corpus. The two adapters matter here. `public_intents` was trained on six intent datasets, including banking77 and clinc150, which are also public-menu tasks. `typed_decisions` was fitted with labels on the training split of the public typed-decisions dataset, and its gate opens on every one of the 8,000 typed-decisions rows we use. That training split is exactly the pool our feedback methods learn from on typed-decisions. So on that benchmark "base Vega" is a model that has already seen the labels of its feedback rows, and it is 5 to 6 points more accurate on the pool than on the test split; we report typed-decisions on Vega both as released and with the shipped adapters switched off (Section 6.7). Vega and Laya come from the same developer, and we pin vegaml 0.8.0, the version the README's benchmark notebook uses.

Laya. Laya [6] is a ModernBERT-large encoder [53] (395M parameters) whose input carries the situation, the question and every option, each option introduced by a marker token, followed by a two-layer transformer head (26.5M parameters) that scores the marker positions; the scores are divided by a shipped temperature per (question type, option-count bucket). It was trained with RLCD on undisclosed data and reads 512 tokens, so rows whose option markers do not fit are unanswerable for every method built on it (all of clinc150_oos; 49 DecisionBench test rows). `laya-typed-decisions` is the developer's whole-model fine-tune on typed-decisions' 6,000 training decisions, with a 1,024-token window. On the five other benchmarks it is a transfer model that never saw a row, which is how we present it.

4 Peekaboo

4.1 Setting, and what "one bit" means under each protocol

A decision model is a policy $\pi_0(a \mid s)$ over the options a of a typed question s. In deployment it commits to an answer and later receives $r \in \{0, 1\}$. We train an add-on π from a pool of (question, answer, verdict)

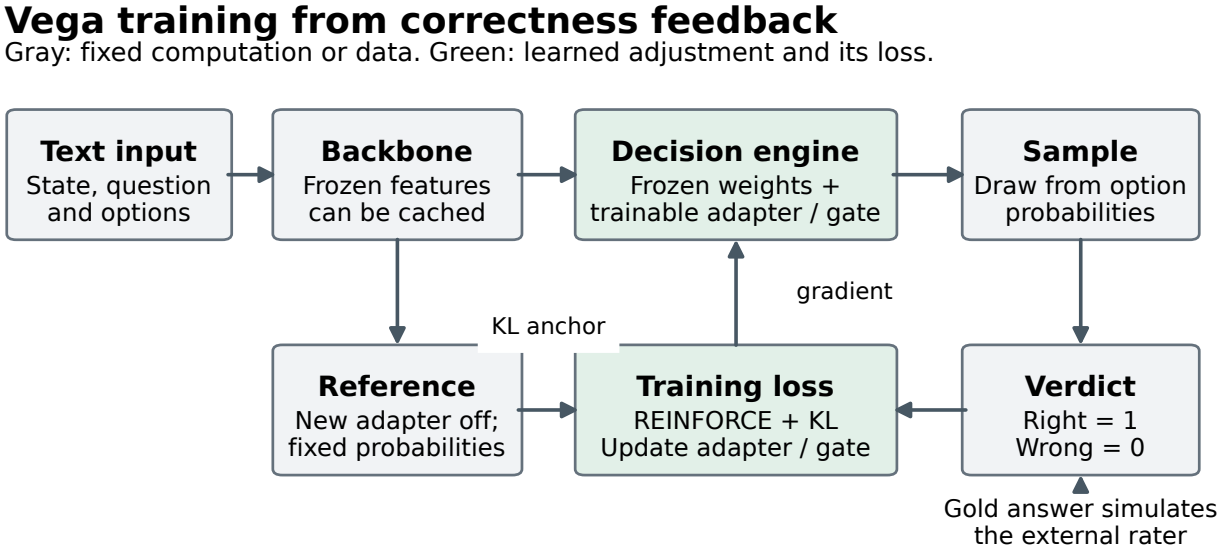


Figure 1: Vega training architecture. A frozen backbone produces cached features. The decision engine retains its original weights while a new adapter and gate learn from sampled-answer correctness. The reference branch disables only the new adapter and supplies the KL anchor. Gold answers simulate the external rater; they are not policy inputs. One pass collects one verdict per training row; eight passes collect fresh verdicts on repeated rows. Temperature fitting and the optional model-level switch use additional held-out feedback.

records and score it on held-out test rows. The reward is computed from the benchmark's gold answer, so it is never wrong or missing unless we make it so (Section 6.1). Two protocols appear throughout:

- **eight passes**: the policy answers every training row once per pass for eight passes, each time sampling a fresh answer from its current probabilities and receiving a fresh verdict. A row can thus be rated up to eight times on different answers. Whenever a verdict is 1 the gold answer has been revealed outright, and for two-option questions any verdict is the label. Section 6.3 measures how often this happens: for most rows, it does. It is the repeated-feedback comparison; the matched labelled reference also makes eight passes.
- **one pass**: every training row is answered once and rated once. This is our primary feedback-budget setting, motivated by a deployment that rates each served answer once. It also takes one eighth of the gradient steps, so we run the answer-key arms for one pass as well.

4.2 The add-on

The released reference weights remain unchanged. Freezing those parameters does not freeze the output of the adapted model: gradients pass through its computations to update the added parameters. On Vega a gated LoRA adapter [22] of rank 8 is attached to seven of the engine's linear layers with VegaML's own `attach_adapter` (409k parameters at 0.8B, 857k at 4B, most of them in the gate network); its gate starts open and the shipped adapters keep routing as released. The policy is Vega's usual calibrated output computed from the adapted energies. On Laya we train a copy of the decision head with the encoder frozen (*head only*),

or a copy of head and encoder (*encoder unfrozen*) at Laya's own encoder learning rate of 2.5×10^{-5}. The second is a fine-tune of the whole 421M model with a KL anchor, not a small add-on, and we describe it as such.

4.3 The update, and what it is in expectation

For a batch of questions the policy samples $a \sim \pi(\cdot \mid s)$, receives $r = \mathbf{1}[a \text{ is right}]$, and minimises

$$\mathcal{L} = -(r - b) \log \pi(a \mid s) + \beta \, \mathrm{KL}\big(\pi(\cdot \mid s) \,\|\, \pi_0(\cdot \mid s)\big), \qquad \beta = 0.1, \tag{1}$$

with b an exponential running mean of recent rewards ($b \leftarrow 0.9b + 0.1\bar{r}$, started at 0.5), Adam at 10^{-3} (Vega) or 10^{-4} (Laya head), batches of 32, gradient norm clipped at 1. Because the baseline is updated only after the batch, the expected gradient of the first term is $-\pi(y \mid s)\nabla \log \pi(y \mid s)$, the supervised cross-entropy gradient weighted by the policy's current probability of the gold answer. The answer-key arm minimises $-\log \pi(y \mid s)$ with the same KL term, rows, epochs and optimiser: the two differ only in that weight and in the sampling noise. That is the right frame for every "feedback against labels" number in this paper. This is an identity of expected gradients; differentiating a probability-weighted cross-entropy loss would add another term. The weighting predicts that feedback will trail labels most on rows the base model ranks the gold answer low, and tie or lead where the gold answer is already near the top; Section 6.3 checks this. The fixed β was chosen in a pilot that trained and tested on DecisionBench's README rows, from $\{0, 0.1, 1\}$; it was not retuned, and the $\beta = 1$ runs of that grid are reported in Appendix A.

4.4 Temperature scaling from correctness verdicts

Policy-gradient training on an improper reward sharpens the probabilities (Section 6.4 gives the pre-fix ECE). We refit one temperature T on a held-out calibration split (20% of each task's feedback rows) by sampling $a \sim \pi$ at $T = 1$, observing r, and choosing

$$T^\star = \arg\max_T \sum_i \big[r_i \log p_T(a_i \mid s_i) + (1 - r_i) \log(1 - p_T(a_i \mid s_i)) \big], \qquad p_T = \mathrm{softmax}(\log p / T), \tag{2}$$

over 61 log-spaced values in $[0.25, 32]$. Since the trained policy samples its argmax most of the time, this lands within one grid step of a fit on argmax correctness in 69 of 82 runs (on the same step in 27); it is top-label temperature scaling [19] from the verdicts a deployment already has. A temperature leaves the argmax unchanged, so accuracy is identical with and without it. The probabilities are floored at 10^{-12} before scaling, which ties the non-argmax options on rows where the trained policy had driven them below the floor; this affects NLL on those rows and is part of the evaluated probability-processing pipeline.

4.5 The switch

On the calibration rows, the add-on's best answer and the released model's best answer are each rated (two verdicts per row), and the add-on is kept only if it was right more often. For comparison, we also rate one *sampled* answer from each model. That statistic estimates sampled-answer correctness, not deployed best-answer accuracy, and can move differently when probabilities are sharpened or temperature-scaled. The sampled statistic is still computed and reported for comparison. It needs two verdicts on the same question, one of them on an answer the system did not serve, so a deployment would run it as a held-out comparison on a slice of traffic. The switch is a safeguard on the add-on as a whole, not a per-decision validator.

Table 1: The six benchmarks. "Feedback rows" are the rows whose verdicts a method may learn from (80%) and calibrate on (20%); no test row is ever used for learning.

benchmark	reported by	test rows	feedback rows	rows and pool
43 public tasks	Vega	2,535	2,134	43 public datasets as typed questions, up to 60 test rows each; feedback: VegaML's 50 fit rows per task, drawn class by class; macro-averaged over tasks
JevBench	Vega	231	(5-fold)	rules, policies and judgment calls in 18 families, paraphrase groups; no training split, every method cross-fits
DecisionBench	Vega	1,075	3,205	43 tasks in 28 domains; the README's 25-row slice per task is the test set, other rows of the same tasks the pool
typed-decisions	Laya	2,000	6,000	400 test and 1,200 training cases, five decisions each; the pool is the public training split, also the training data of Vega's shipped adapter
emotion	(ours)	1,503	4,000	six emotions of a tweet; pool: training rows drawn at random (natural class mix)
AG News	(ours)	1,912	4,000	four news sections; pool: training rows drawn at random

5 Experimental setup

5.1 Benchmarks and pools

Table 1 and Figure 2 describe the six benchmarks: the three in Vega's README (43 public tasks, JevBench [49], DecisionBench [21]), the one in Laya's README (typed-decisions [35]), and two classification sets we added for their large training splits (emotion [46], AG News [58]). Rows are read, asked and scored with VegaML's benchmark harness, vendored unmodified, and the test rows are fixed across the methods compared here. Three things about the pools are disclosed here. The public menu's feedback pool is the 50 "fit" rows per task that VegaML's notebook gives its TTT head, drawn class by class, so it is class-balanced while the 60 test rows per task are not; a class-balanced pool can confound adaptation when test class frequencies differ. The emotion and AG News pools are 4,000 rows drawn at random from the whole training split, so they retain the dataset's natural class mix; correspondence with deployment traffic is not measured. The DecisionBench pool is the other rows of the README's 43 tasks (100 per task minus the 25-row test slice), drawn at random. On typed-decisions the pool is the public training split, which is also the training data of Vega's shipped adapter (Section 3). Two public tasks and three DecisionBench tasks have construction artifacts that an adapter can learn without reading the input: in `truthfulqa_mc1` the harness keeps the upstream option order, so the gold answer is option A in every row; in `pubmedqa` the gold is the majority class on 55% of test rows; in the two DecisionBench tasks built from Civil Comments (`contains_threat` and `toxicity_severity`) the gold is the first option in 86% of rows; in `finqa-numerical-reasoning` the options are always $\{x, 100x, x/100, 1.1x\}$ with gold x. Table 9 reports the main public-menu and DecisionBench gains with and without these tasks.

5.2 Protocol and statistics

Every method learns from its pool's feedback rows (or their labels) and is scored on the test rows. On JevBench, which has no training split, every method including the answer-key arms learns by 5-fold cross-fitting with all paraphrases of a question kept in one fold, and the calibration fifth of each training fold is drawn by paraphrase group. The feedback rows of each task are split 80 / 20 into training and calibration rows with a

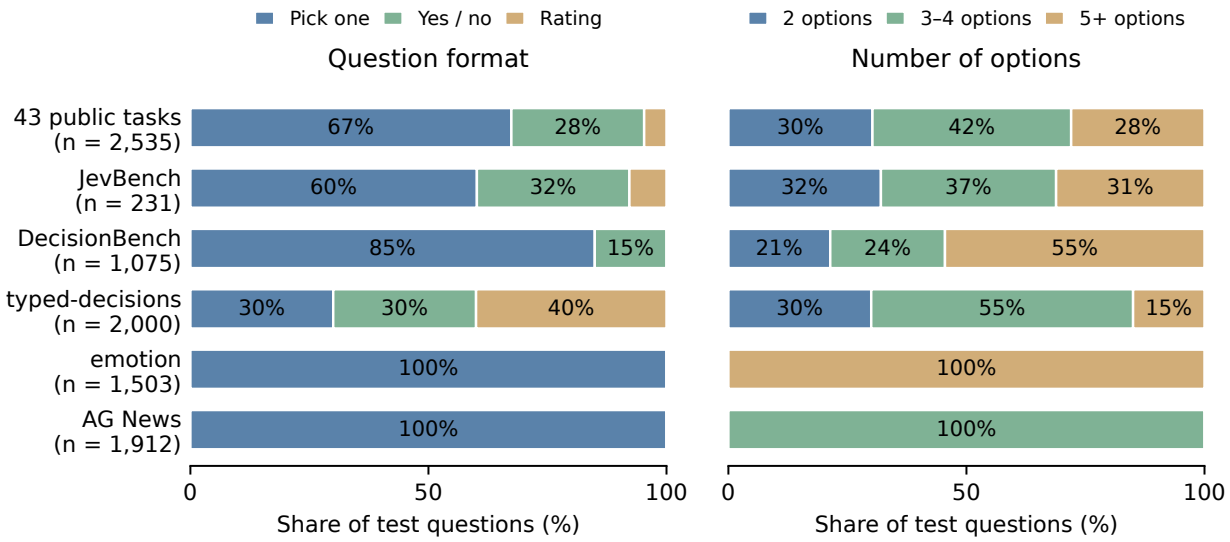


Figure 2: Benchmark composition. Each row gives the benchmark's test count and percentages by answer format (left) and option count (right). Both panels use a 0–100% scale. Two-option questions matter because either correctness verdict identifies the label.

fixed permutation. The main Vega feedback and answer-key arms use seeds 0, 1 and 2. Other arms use fewer seeds as listed in Table 12; no training-seed uncertainty can be estimated from a single run. Gains are paired differences in per-row correctness between two methods on the same rows, in points. Their 95% intervals come from a bootstrap that resamples the training seeds of each side with replacement and then resamples rows by cluster: tasks on the public menu (which is macro-averaged over tasks) and DecisionBench, cases on typed-decisions, paraphrase groups on JevBench. With at most three seeds, estimates of training variation are imprecise; single-seed intervals are conditional on that run and resample test rows or clusters only. Two-sided p-values come from the same bootstrap and are Holm-adjusted within each family of comparisons a reader is likely to count (a method against the base model across benchmarks; Peekaboo against each answer-key arm across benchmarks; the KL contrasts). "Non-inferior" marks an unadjusted 95% interval whose lower bound is above -2 points; these exploratory marks are not multiplicity-adjusted. Failure to reject a difference does not establish equality or non-inferiority. Same-seed reruns on different machines move individual results by up to 3 points (the earlier grid's typed-decisions seed 2 moved by 3.1), which limits the precision of cross-machine comparisons even with fixed seeds. Accuracy, ECE (15 equal-width bins, as in the README), NLL, Brier, AURC, macro-F1 and the per-benchmark slices come from the harness's `score_metrics`.

Feedback accounting. "One rating per row" refers to the training subset, not the complete adaptation pipeline. Temperature fitting uses a sampled-answer verdict on each held-out calibration row. The optional switch additionally compares two best-answer verdicts per calibration row; the sampled-switch diagnostic collects further verdicts for analysis. The early-stopped supervised reference uses calibration labels to choose an epoch. Our comparisons match training rows and report these additional information sources, but do not establish equal total annotation effort across pipelines or the cost of collecting human verdicts.

5.3 Methods compared

- **Released model**: Vega 0.8B, Vega 4B, Laya, replayed on cached inputs and checked against the official prediction call and the README numbers. On typed-decisions, Vega also with its shipped adapters off.

- **Their own recipe**: for Vega, VegaML's TTT head trained on the labelled feedback rows of each task, under the vegaml 0.8.0 rule (a head replaces the engine where its cross-validated accuracy beats chance) and the 0.12.3 rule (where it also beats the engine on the same rows); for Laya, the released `laya-typed-decisions` checkpoint on every benchmark, in-domain only on typed-decisions.
- **Answer-key add-on, same recipe**: the same adapter (Vega) or head copy (Laya) trained with cross-entropy on the labels of the same rows, the same KL term, eight passes, then label-based temperature scaling; also run for one pass.
- **Answer-key add-on, early-stopped**: the same run with the epoch chosen by argmax accuracy on the calibration rows' labels.
- **Answer-key fine-tune, encoder unfrozen** (Laya): head and encoder trained with cross-entropy, the same KL term, rows and passes as the encoder-unfrozen Peekaboo; with and without early stopping.
- **Label prior**: one log-offset per (task, option text) on the released probabilities, no input features, trained with Peekaboo's own loop from verdicts (learning rate chosen on the calibration rows' verdicts) and, separately, from the labels.
- **Released model + temperature**: the released model with the honesty fix applied to its own sampled answers on its calibration rows, and with a label-fitted temperature.
- **Peekaboo** with eight passes and with one pass; head only and encoder unfrozen on Laya; without the KL term; with the switch.
- **Realistic feedback**: the released model's best answers rated once and learned off-policy (a verdict of 1 is a label, a verdict of 0 rules that answer out); its sampled answers rated once and learned with clipped importance weights; eight-pass Peekaboo with 20% of verdicts flipped, with only half of the answers rated, and with every wrong answer reported but only 20% of right ones.

5.4 Compute and reproducibility

The grid (314 steps: caches, replays and every method and seed) ran in about ten hours on one NVIDIA A40 from cached backbone features; the seed counts that budget allowed are listed in Section 7. One script, `experiments/peekaboo_remaining/run.py`, runs every step, skips finished ones and retries failures; `analysis.py` computes everything in this paper from the per-row records; `paper/make_figures.py` writes every table, figure and number in the text, to reduce transcription errors; interpretation still requires manual checking against those records. Historical results remain under `experiments/peekaboo_remaining/output/`; they are not pooled with the main grid.

6 Results

6.1 One rating per row, and feedback as a deployment delivers it

Figure 3 shows what a single pass keeps. On DecisionBench one-pass Peekaboo gains +6.7 [+3.6, +10.5] at 0.8B and +3.2 [+1.5, +5.4] at 4B, against +11.3 and +5.7 with eight passes; the one-pass labelled adapter gains +10.0 and +4.7. At matched epochs the gap between verdicts and labels is −3.2 [−5.8, −0.6] and −1.4 [−3.8, +0.8], wider than with eight passes. On typed-decisions one pass gives +4.1 and +2.0 (labels: +6.8, +4.6; verdicts significantly behind at both sizes); on the public menu +1.9 and +0.0; on emotion +5.0 and +3.2; on AG News +1.3 and +0.2. One pass keeps about half to two thirds of the eight-pass gain on DecisionBench and the public menu and more on emotion.

One pass confounds two things: each row is rated once, and the model takes one eighth of the gradient steps. The logged arms allow additional optimisation without new training verdicts, but also change the collection policy and use an off-policy objective. Each training row is answered once by the released model, rated once, and the add-on is then trained on that fixed log for eight passes, off-policy. With the released

Table 2: Accuracy, mean over seeds (± s.d.; three seeds except where Section 7 lists fewer). JevBench has no training split: every method learns by 5-fold cross-fitting there.

model	method	learns from	public	JevBench	DecisionB.	typed-dec.	emotion	AG News
Vega 0.8B	released model	none	0.556	0.545	0.509	0.666	0.568	0.901
	TTT head, 0.8.0 rule	answer key	0.516	·	0.530	0.719	0.691	0.900
	TTT head, 0.12.3 rule	answer key	0.576	·	0.532	0.719	0.691	0.901
	answer-key adapter, same recipe	answer key	0.590 ±0.004	0.537 ±0.020	0.632 ±0.004	0.737 ±0.002	0.679 ±0.007	0.901 ±0.003
	answer-key adapter, early-stopped	answer key	0.603 ±0.010	0.538 ±0.005	0.641 ±0.010	0.740 ±0.003	0.676 ±0.006	0.917 ±0.002
	answer-key adapter, one pass	answer key	0.592 ±0.003	0.540 ±0.002	0.608 ±0.007	0.734 ±0.007	0.635 ±0.002	0.917 ±0.002
	Peekaboo, eight passes	verdicts	0.593 ±0.013	0.522 ±0.010	0.622 ±0.002	0.728 ±0.007	0.627 ±0.004	0.912 ±0.004
	Peekaboo, one pass	verdicts	0.575 ±0.008	0.540 ±0.002	0.576 ±0.011	0.708 ±0.003	0.618 ±0.006	0.914 ±0.002
	logged best answers, rated once	verdicts	·	·	0.583 ±0.009	0.726	0.603	·
	logged sampled answers, rated once	verdicts	·	·	0.592 ±0.011	0.722	0.634	·
	Peekaboo, no KL term	verdicts	0.582	0.517 ±0.014	0.609 ±0.004	0.709	0.622	0.916
Vega 4B	released model	none	0.718	0.667	0.636	0.729	0.575	0.928
	TTT head, 0.8.0 rule	answer key	0.678	·	0.640	0.757	0.697	0.915
	TTT head, 0.12.3 rule	answer key	0.724	·	0.654	0.755	0.697	0.915
	answer-key adapter, same recipe	answer key	0.706 ±0.002	0.706 ±0.011	0.687 ±0.007	0.769 ±0.001	0.640 ±0.001	0.924 ±0.003
	answer-key adapter, early-stopped	answer key	0.726 ±0.005	0.677 ±0.010	0.688 ±0.005	0.760 ±0.008	0.663 ±0.004	0.931 ±0.004
	answer-key adapter, one pass	answer key	0.723 ±0.003	0.667 ±0.000	0.683 ±0.009	0.775 ±0.008	0.631 ±0.004	0.931 ±0.003
	Peekaboo, eight passes	verdicts	0.706 ±0.005	0.696 ±0.005	0.694 ±0.005	0.746 ±0.010	0.629 ±0.009	0.930 ±0.002
	Peekaboo, one pass	verdicts	0.718 ±0.017	0.664 ±0.003	0.669 ±0.004	0.749 ±0.004	0.607 ±0.003	0.930 ±0.002
	logged best answers, rated once	verdicts	·	·	0.655 ±0.004	0.769	0.611	·
	logged sampled answers, rated once	verdicts	·	·	0.680 ±0.003	0.748	0.620	·
	Peekaboo, no KL term	verdicts	0.713	0.693	0.672 ±0.002	0.742	0.519	0.928
Laya	released model	none	0.540	0.576	0.406	0.362	0.599	0.930
	laya-typed-decisions (fine-tune on typed-decisions)	answer key	0.550	0.541	0.436	0.761	0.604	0.930
	answer-key head, one pass	answer key	0.541	0.576	0.403	0.412 ±0.011	·	·
	Peekaboo, head only, one pass	verdicts	0.541	0.576	0.406	0.364 ±0.003	·	·
	answer-key fine-tune, encoder unfrozen	answer key	0.587 ±0.009	0.535 ±0.025	0.560 ±0.016	0.750	0.899	·
	answer-key fine-tune, encoder unfrozen, early-stopped	answer key	0.584 ±0.003	0.595 ±0.005	0.554 ±0.005	0.750	0.905	·
	Peekaboo, encoder unfrozen, eight passes	verdicts	0.566 ±0.001	0.558 ±0.030	0.545 ±0.014	0.685	0.820	·
	Peekaboo, encoder unfrozen, one pass	verdicts	0.562 ±0.005	0.573 ±0.005	0.482 ±0.026	0.588	0.698	·

Table 3: Feedback as a deployment delivers it: gain over the released model in points (95% intervals), on the three benchmarks with the realistic-feedback arms.

feedback	DecisionBench			typed-decisions			emotion		
	Vega 0.8B	Vega 4B	Laya (head)	Vega 0.8B	Vega 4B	Laya (head)	Vega 0.8B	Vega 4B	Laya (head)
eight passes, clean verdicts	+11.3 [+7.0, +16.1]	+5.7 [+2.7, +9.3]	–	+6.2 [+4.2, +8.2]	+1.7 [-0.1, +3.6]	–	+5.9 [+3.9, +8.0]	+5.4 [+3.4, +7.4]	–
one pass, clean verdicts	+6.7 [+3.6, +10.5]	+3.2 [+1.5, +5.4]	+0.1 [+0.0, +0.3]	+4.1 [+2.2, +6.1]	+2.0 [+0.7, +3.5]	+0.2 [-0.2, +0.5]	+5.0 [+3.1, +6.9]	+3.2 [+1.6, +4.8]	–
logged: the released model's best answers, rated once	+7.4 [+3.4, +12.2]	+1.8 [-1.4, +5.3]	–	+6.0 [+3.9, +8.2]	+4.0 [+2.5, +5.7]	–	+3.5 [+0.7, +6.2]	+3.6 [+1.2, +5.9]	–
logged: the released model's sampled answers, rated once (clipped IPS)	+8.3 [+4.4, +13.0]	+4.3 [+1.6, +7.3]	–	+5.6 [+3.5, +7.6]	+1.8 [+0.1, +3.8]	–	+6.6 [+4.6, +8.6]	+4.5 [+2.5, +6.6]	–
20% of verdicts flipped (eight passes)	+9.1 [+5.8, +12.8]	+4.2 [+1.7, +7.3]	–	+5.9 [+4.2, +7.5]	+2.8 [+1.2, +4.3]	–	+4.7 [+3.1, +6.3]	+4.7 [+2.7, +6.8]	–
half of the answers rated (eight passes)	+10.3 [+6.2, +14.8]	+4.7 [+1.6, +8.3]	–	+5.2 [+3.1, +7.5]	+3.5 [+1.8, +5.3]	–	+6.5 [+4.2, +8.9]	+0.0 [+0.0, +0.0]	–
every wrong answer reported, 20% of right ones (eight passes)	+10.0 [+6.0, +14.7]	+5.7 [+2.7, +9.1]	–	+7.0 [+5.0, +9.0]	+4.2 [+2.4, +5.9]	–	+6.7 [+4.5, +9.0]	+2.7 [+0.9, +4.5]	–
answer key, one pass	+10.0 [+6.3, +13.9]	+4.7 [+2.1, +7.4]	-0.3 [-0.7, +0.0]	+6.8 [+4.8, +8.8]	+4.6 [+3.0, +6.4]	+5.0 [+1.9, +7.7]	+6.6 [+4.5, +8.8]	+5.7 [+3.8, +7.7]	–
answer key, eight passes	+12.3 [+7.8, +17.5]	+5.1 [+1.4, +9.1]	–	+7.0 [+5.0, +9.1]	+4.0 [+2.5, +5.6]	–	+11.1 [+8.6, +13.6]	+6.5 [+4.3, +8.8]	–

model's sampled answers and clipped importance weights this gains +8.3 [+4.4, +13.0] on DecisionBench at 0.8B and +4.3 [+1.6, +7.3] at 4B. Against one on-policy pass that is +1.6 [−1.4, +4.8] and +1.1 [−1.2, +3.9]; against eight passes it is −2.9 [−5.3, −0.9] and −1.4 [−2.9, +0.1]. On DecisionBench the logged arms have three seeds. Their point estimates recover part of the one-pass shortfall while remaining below eight fresh-feedback passes; the latter difference survives Holm adjustment at 0.8B. This is consistent with a benefit from additional verdicts, but is not a controlled isolation of information from policy and objective changes. On typed-decisions and emotion, where the logged arms have one seed, the logged-versus-eight-pass differences are not significant (−0.6 and +0.7 at 0.8B, +0.2 and −0.9 at 4B). Against the early-stopped labelled adapter the logged sampled arm is −4.8 [−7.7, −2.0] on DecisionBench, −1.8 on typed-decisions and −4.2 on emotion at 0.8B, and −0.8, −1.2 and −4.3 at 4B: one logged verdict per row stays below the labels on every cell, significantly on four of the six.

Table 3 and Figure 4 collect the deployment question. A system that serves its best answer rather than a sample logs only that answer's verdict: a 1 is a label, a 0 rules one option out. Learning from that log gains +7.4 [+3.4, +12.2] on DecisionBench at 0.8B and +1.8 [−1.4, +5.3] at 4B, +6.0 and +4.0 on typed-decisions,

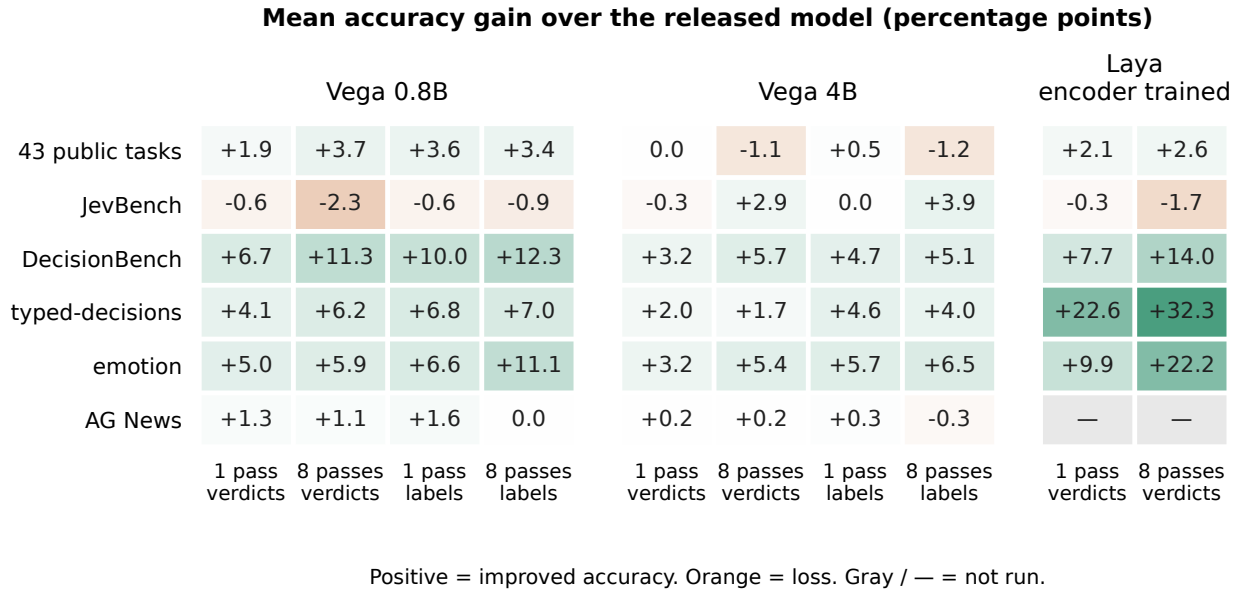


Figure 3: Feedback budget and optimisation passes. Each cell is mean accuracy gain over the released model, in percentage points; positive values indicate improvement. Columns identify one or eight passes and verdicts or labels. Labelled columns use the same-recipe arm without early stopping. Laya trains the encoder and head. Gray cells were not run. These are point estimates; Tables 4 and 5 provide uncertainty. One pass collects one verdict per training row and uses one eighth of the updates; calibration feedback is additional.

and +3.5 and +3.6 on emotion. It trails the early-stopped labels by 5.8 and 3.3 points on DecisionBench and by 7.3 and 5.3 on emotion, and on typed-decisions at 4B the difference from them is not significant (+1.0). Neither logging policy consistently dominates; the point-estimate advantages of sampled over best-answer logging are +0.9 and +2.5 points more on DecisionBench and +3.1 on emotion at 0.8B, and −2.2 on typed-decisions at 4B.

We measured damaged feedback only under the eight-pass protocol, where a flipped or missing verdict on one pass is redrawn on the next; how noise affects a single rating per row is not measured here. On DecisionBench, with three seeds, 20% of verdicts flipped costs 2.2 points at 0.8B and 1.5 at 4B, half of the answers unrated 1.0 and 1.0, and complaint-skewed verdicts 1.2 and 0.0. On typed-decisions and emotion these arms have one seed against the clean arm's three and differ from it by up to three points in either direction. One run failed: with half of the answers unrated on emotion at 4B the adapter's predictions are identical to the released model's on every test row and its reward stays flat through all eight passes, so that cell (+0.0) says nothing about missing feedback. Rewards are computed from the gold answer in all of these; a rater who is systematically wrong about what was right is not simulated beyond the flipped verdicts.

6.2 Eight passes: verdicts against labels, and against the vendors' own recipes

Tables 2 and 4 give accuracy and gain over the released model for every arm, Table 5 the paired head-to-heads, and Figure 5 the direct differences against labelled comparators. Under the eight-pass protocol the verdict-trained add-on lands close to the same add-on trained on labels and early-stopped, and mostly below it. On DecisionBench Peekaboo gains +11.3 points at 0.8B and +5.7 at 4B; the early-stopped labelled adapter gains +13.2 and +5.1. The paired difference is −1.9 [−4.1, +0.3] at 0.8B and +0.6 [−1.2, +2.3] at 4B; on typed-decisions −1.3 and −1.4, on the public menu −1.1 and −2.0, on AG News −0.5 and −0.1, on JevBench −1.6 and +1.9. Emotion is the exception: −4.9 [−6.8, −3.0] and −3.5 [−5.3, −1.7], and these two are the

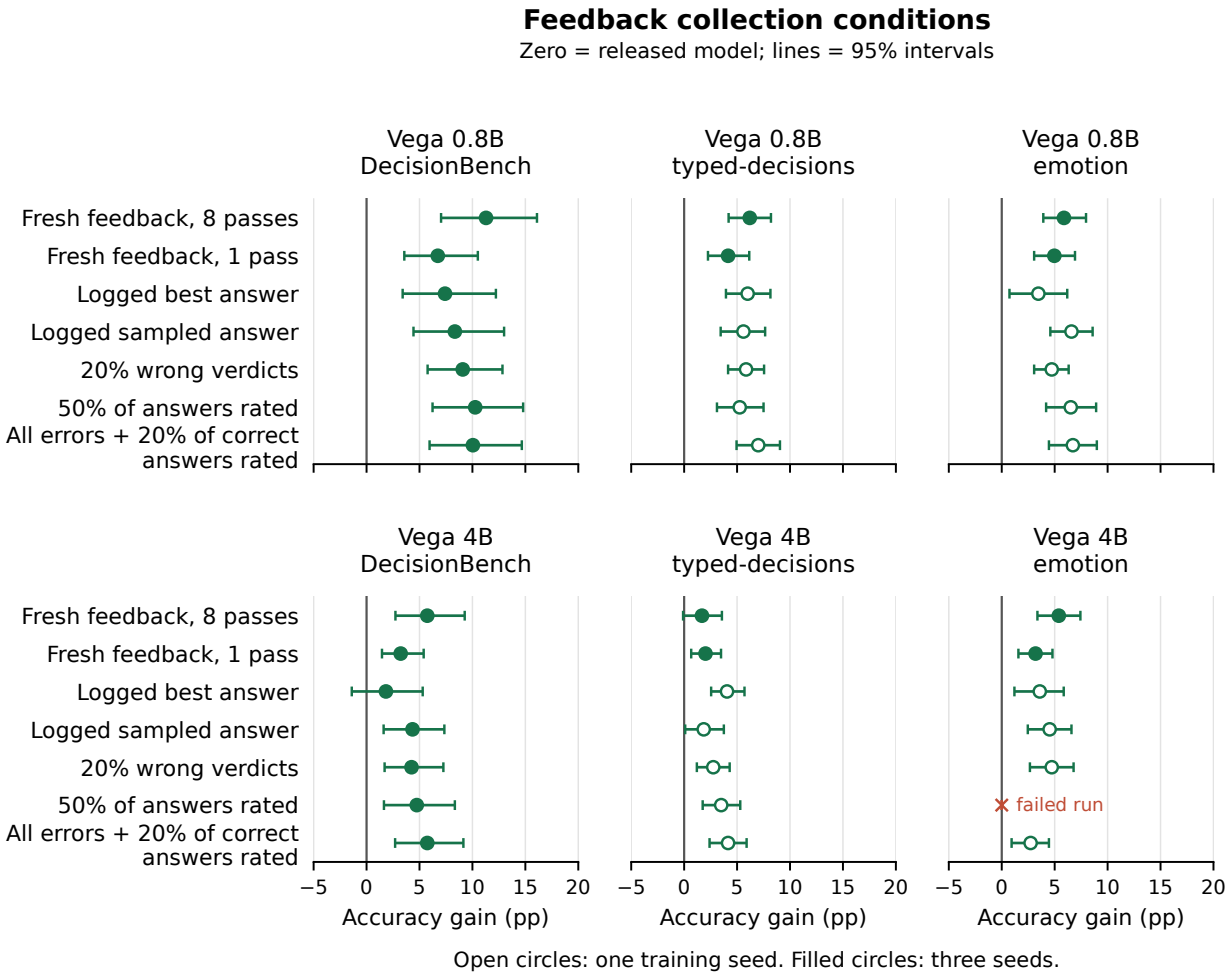


Figure 4: Feedback collection conditions, separated by model and benchmark. Each point is accuracy gain over the released model and each line its 95% interval. Clean arms use one or eight fresh-feedback passes; logged arms reuse one verdict per training row for eight optimisation passes; damaged-feedback arms use eight passes. Filled circles use three training seeds; open circles use one, whose intervals cannot estimate training variability. The red cross identifies the failed half-rated emotion run at 4B (Section 6.1); it is not evidence that missing feedback produces zero gain.

only differences in the family that survive Holm adjustment. Against the same-recipe adapter without early stopping the differences are smaller still. Repeated label revelation (Section 6.3) gives the feedback arm access to much of the same information as supervised training, although its updates retain different weighting and sampling noise. These are point-estimate comparisons, not a demonstration of equivalence: repeated positive verdicts identify labels for many rows. The one-rating arms in Section 6.1 measure a single training verdict per row.

Against VegaML's TTT head. Under the evaluated 0.12.3 rule, which keeps a head only where it beats the engine on its own fit rows, the TTT head gains +2.0 and +0.6 on the public menu; under the 0.8.0 rule it lost 4.1 and 3.9 points, showing that the choice of selection rule materially changes the baseline. Peekaboo minus the guarded TTT head is +9.0 [+5.4, +12.6] on DecisionBench at 0.8B and +4.0 [+1.5, +6.5] at 4B; +1.1 [+0.1, +2.0] and +1.5 [+0.4, +2.6] on AG News; +1.7 [−1.6, +5.0] and −1.8 [−4.3, +0.7] on the public menu; +0.9 and −0.9 on typed-decisions; and −6.4 [−8.4, −4.3] and −6.8 [−8.9, −4.7] on emotion. After

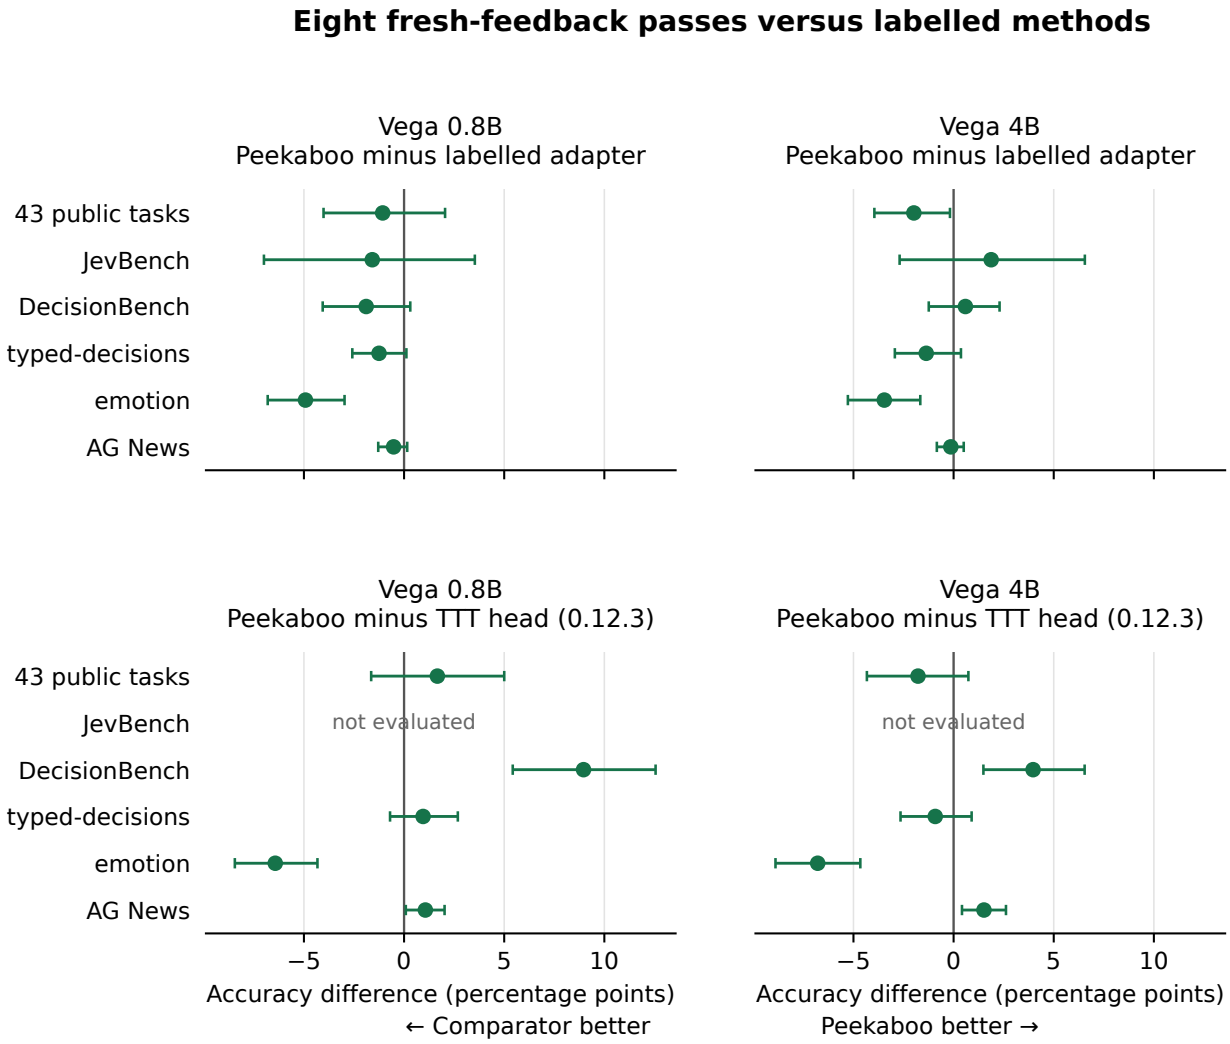


Figure 5: Direct comparison of eight-pass Peekaboo with labelled adaptation on Vega. Points are Peekaboo accuracy minus comparator accuracy, in percentage points, with 95% intervals. Left of zero favors the comparator; right favors Peekaboo. Top panels use the matched labelled adapter with early stopping; bottom panels use the vendor TTT head under the evaluated 0.12.3 rule. The unavailable TTT comparison on JevBench is labelled. Intervals crossing zero are inconclusive; Holm-adjusted decisions appear in Table 5. Laya is shown separately in Figure 10.

Holm adjustment the verdict-trained adapter is ahead of VegaML's labelled recipe on DecisionBench at both sizes and on AG News at 4B, without a significant difference on the public menu and typed-decisions, and behind on emotion at both sizes. On DecisionBench the guarded head answers 329 of 1,075 rows at 0.8B and 302 at 4B (Table 8). On the rows it answers, Peekaboo minus TTT is +3.2 [−2.0, +7.9] and +3.1 [−1.9, +8.9], neither significant, and the labelled adapter does about as well there (+3.9, +3.6). The significant margin is on the rows the head leaves to the engine (+11.5 and +4.3), where it is a comparison with the released model, and it belongs to an in-engine adapter that reads the option text, not to the feedback signal. Emotion goes to TTT by a wide margin: a six-way single-label task with 3,200 labelled rows is where a logistic head on good features is at its best.

Table 4: Gain over the released model in accuracy points; † marks differences that survive Holm adjustment within the method's family. The 95% intervals are in Figure 5 and Table 5. JevBench has no training split: every method learns by 5-fold cross-fitting there.

model	method	learns from	public	JevBench	DecisionB.	typed-dec.	emotion	AG News
Vega 0.8B	released model	none	–	–	–	–	–	–
	TTT head, 0.8.0 rule	answer key	-4.1	·	+2.1	+5.3†	+12.3†	-0.2
	TTT head, 0.12.3 rule	answer key	+2.0	·	+2.3	+5.2†	+12.3†	+0.0
	answer-key adapter, same recipe	answer key	+3.4	-0.9	+12.3†	+7.0†	+11.1†	+0.0
	answer-key adapter, early-stopped	answer key	+4.7†	-0.7	+13.2†	+7.4†	+10.8†	+1.6†
	answer-key adapter, one pass	answer key	+3.6	-0.6	+10.0†	+6.8†	+6.6†	+1.6†
	Peekaboo, eight passes	verdicts	+3.7	-2.3	+11.3†	+6.2†	+5.9†	+1.1
	Peekaboo, one pass	verdicts	+1.9	-0.6	+6.7†	+4.1†	+5.0†	+1.3†
	logged best answers, rated once	verdicts	·	·	+7.4†	+6.0†	+3.5†	·
	logged sampled answers, rated once	verdicts	·	·	+8.3†	+5.6†	+6.6†	·
	Peekaboo, no KL term	verdicts	+2.5	-2.9	+10.0	+4.2	+5.4	+1.5
Vega 4B	released model	none	–	–	–	–	–	–
	TTT head, 0.8.0 rule	answer key	-3.9	·	+0.4	+2.8†	+12.2†	-1.3
	TTT head, 0.12.3 rule	answer key	+0.6	·	+1.8	+2.6†	+12.2†	-1.3
	answer-key adapter, same recipe	answer key	-1.2	+3.9	+5.1†	+4.0†	+6.5†	-0.3
	answer-key adapter, early-stopped	answer key	+0.8	+1.0	+5.1†	+3.0†	+8.8†	+0.3
	answer-key adapter, one pass	answer key	+0.5	-0.0	+4.7†	+4.6†	+5.7†	+0.3
	Peekaboo, eight passes	verdicts	-1.1	+2.9	+5.7†	+1.7	+5.4†	+0.2
	Peekaboo, one pass	verdicts	+0.0	-0.3	+3.2†	+2.0†	+3.2†	+0.2
	logged best answers, rated once	verdicts	·	·	+1.8	+4.0†	+3.6†	·
	logged sampled answers, rated once	verdicts	·	·	+4.3†	+1.8†	+4.5†	·
	Peekaboo, no KL term	verdicts	-0.5	+2.6	+3.5	+1.3	-5.6	+0.0
Laya	released model	none	–	–	–	–	–	–
	laya-typed-decisions (fine-tune on typed-decisions)	answer key	+1.0	-3.5	+3.1	+40.0†	+0.5	+0.0
	answer-key head, one pass	answer key	+0.1	+0.0	-0.3	+5.0†	·	·
	Peekaboo, head only, one pass	verdicts	+0.1	+0.0	+0.1	+0.2	·	·
	answer-key fine-tune, encoder unfrozen	answer key	+4.6†	-4.0	+15.4†	+38.8†	+30.0†	·
	answer-key fine-tune, encoder unfrozen, early-stopped	answer key	+4.3	+1.9	+14.8†	+38.8†	+30.6†	·
	Peekaboo, encoder unfrozen, eight passes	verdicts	+2.6	-1.7	+14.0†	+32.3†	+22.2†	·
	Peekaboo, encoder unfrozen, one pass	verdicts	+2.1	-0.3	+7.7†	+22.6†	+9.9†	·

JevBench. On JevBench both feedback and answer-key adaptation use cross-fitting. At 4B Peekaboo gains +2.9 $[-2.5, +8.4]$ and the labelled adapter +3.9 $[-2.3, +10.0]$; at 0.8B Peekaboo loses 2.3 points and the labelled adapter 0.9. None of these intervals excludes zero. JevBench has 231 rows and 38-to-42-row calibration folds.

6.3 What the gain is made of

How much the protocol reveals. Table 6 (right) counts what eight on-policy passes tell the learner. On DecisionBench at 0.8B each training row produces 4.88 reward-1 events on average, each of which names the gold answer; from the released model's own probabilities, 8 draws would hit the gold at least once on a share 0.76 of the rows, and 0.21 of the rows have two options, where any verdict is the label. On typed-decisions the figures are 6.12 reveals per row and 0.94; on emotion 4.90 and 0.85. The first pass's reward, measured before each training row is revisited, is 0.48 on DecisionBench against the released model's best-answer accuracy of 0.50 on the same rows: this is a descriptive sampled-versus-best-answer comparison on the feedback pool. It does not bound cumulative exploration cost or subsequent deployment regressions, and earlier batches may already have updated the policy before a later row is first sampled.

A label prior with no input. Per-task log-offsets on the released probabilities, one per (task, option text), trained with Peekaboo's own loop from verdicts, gain +4.3 on typed-decisions at 0.8B and +1.8 at 4B, against

Table 5: Head to head, in accuracy points with 95% intervals (seeds and row clusters resampled). Positive means the first method is better. †: survives Holm adjustment within the row's family; NI: exploratory non-inferiority at a 2-point margin using an unadjusted interval (lower bound above −2); these marks are not multiplicity-adjusted. TTT is scored under the vegaml 0.12.3 rule; for Laya "their method" is the typed-decisions fine-tune, in-domain only on typed-decisions. The Laya head-only one-pass arms barely move from the released model (Section 6.6), so no NI marks are given for them.

model	comparison	public	JevBench	DecisionB.	typed-dec.	emotion	AG News
Vega 0.8B							
	Peekaboo (8 passes) vs. their method	+1.7 [-1.6, +5.0]	–	+9.0 [+5.4, +12.6]†	+0.9 [-0.7, +2.7]	-6.4 [-8.4, -4.3]†	+1.1 [+0.1, +2.0]
	Peekaboo (8) vs. answer key, same recipe (8)	+0.3 [-2.8, +3.6]	-1.4 [-6.1, +3.0]	-1.0 [-3.0, +1.0]	-0.9 [-2.2, +0.5]	-5.2 [-7.1, -3.3]†	+1.0 [-0.1, +2.1] NI
	Peekaboo (8) vs. answer key, early-stopped	-1.1 [-4.0, +2.0]	-1.6 [-7.0, +3.5]	-1.9 [-4.1, +0.3]	-1.3 [-2.6, +0.1]	-4.9 [-6.8, -3.0]†	-0.5 [-1.3, +0.2] NI
	Peekaboo (1 pass) vs. answer key (1 pass)	-1.7 [-3.7, +0.5]	-0.0 [-1.0, +1.0] NI	-3.2 [-5.8, -0.6]	-2.7 [-4.1, -1.2]†	-1.7 [-3.0, -0.4]†	-0.3 [-0.9, +0.4] NI
	Peekaboo, 1 pass vs. 8 passes	-1.8 [-4.4, +0.8]	+1.7 [-3.2, +7.2]	-4.6 [-8.1, -1.4]†	-2.1 [-3.6, -0.4]	-0.9 [-2.4, +0.5]	+0.2 [-0.5, +1.0]
	logged sampled answers vs. Peekaboo (1 pass)	–	–	+1.6 [-1.4, +4.8] NI	+1.5 [-0.0, +2.9] NI	+1.6 [-0.1, +3.2] NI	–
	logged sampled answers vs. Peekaboo (8)	–	–	-2.9 [-5.3, -0.9]†	-0.6 [-2.1, +0.9]	+0.7 [-1.0, +2.4] NI	–
	logged sampled answers vs. answer key, early-stopped	–	–	-4.8 [-7.7, -2.0]†	-1.8 [-3.3, -0.4]†	-4.2 [-6.1, -2.2]†	–
	logged best answers vs. answer key, early-stopped	–	–	-5.8 [-8.4, -3.3]†	-1.4 [-2.7, -0.1]†	-7.3 [-9.5, -5.3]†	–
	KL term vs. none	+1.1 [-1.5, +3.9]	+0.6 [-2.4, +3.7]	+1.2 [-0.6, +3.0]	+1.9 [+0.3, +3.5]	+0.5 [-1.0, +1.8]	-0.5 [-1.3, +0.3]
	with the switch vs. without	+0.0 [-2.2, +2.3]	+2.3 [-2.9, +7.7]	+0.0 [-1.2, +1.1]	+0.0 [-1.2, +1.3]	+0.0 [-1.1, +1.1]	+0.0 [-0.7, +0.7]
Vega 4B							
	Peekaboo (8 passes) vs. their method	-1.8 [-4.3, +0.7]	–	+4.0 [+1.5, +6.5]†	-0.9 [-2.6, +0.9]	-6.8 [-8.9, -4.7]†	+1.5 [+0.4, +2.6]†
	Peekaboo (8) vs. answer key, same recipe (8)	+0.0 [-2.0, +2.0]	-1.0 [-5.0, +2.8]	+0.6 [-1.6, +2.7] NI	-2.3 [-3.8, -0.8]†	-1.1 [-3.2, +1.0]	+0.6 [-0.1, +1.3] NI
	Peekaboo (8) vs. answer key, early-stopped	-2.0 [-4.0, -0.2]	+1.9 [-2.7, +6.6]	+0.6 [-1.2, +2.3] NI	-1.4 [-2.9, +0.4]	-3.5 [-5.3, -1.7]†	-0.1 [-0.8, +0.5] NI
	Peekaboo (1 pass) vs. answer key (1 pass)	-0.5 [-3.0, +1.4]	-0.3 [-1.4, +0.7] NI	-1.4 [-3.8, +0.8]	-2.6 [-4.0, -1.2]†	-2.5 [-3.7, -1.2]†	-0.1 [-0.7, +0.5] NI
	Peekaboo, 1 pass vs. 8 passes	+1.2 [-1.3, +3.5]	-3.2 [-8.7, +2.4]	-2.5 [-5.8, +0.0]	+0.3 [-1.2, +1.9]	-2.2 [-3.8, -0.5]	+0.0 [-0.6, +0.5]
	logged sampled answers vs. Peekaboo (1 pass)	–	–	+1.1 [-1.2, +3.9] NI	-0.2 [-1.6, +1.4] NI	+1.3 [-0.5, +3.1] NI	–
	logged sampled answers vs. Peekaboo (8)	–	–	-1.4 [-2.9, +0.1]	+0.2 [-1.5, +2.0] NI	-0.9 [-2.7, +0.9]	–
	logged sampled answers vs. answer key, early-stopped	–	–	-0.8 [-2.8, +1.0]	-1.2 [-2.6, +0.4]	-4.3 [-6.4, -2.5]†	–
	logged best answers vs. answer key, early-stopped	–	–	-3.3 [-5.5, -1.2]†	+1.0 [-0.3, +2.4] NI	-5.3 [-7.3, -3.3]†	–
	KL term vs. none	-0.6 [-2.8, +1.4]	+0.3 [-3.1, +3.8]	+2.2 [+0.4, +4.2]	+0.4 [-1.5, +2.3]	+11.0 [+8.5, +13.5]†	+0.2 [-0.2, +0.8]
	with the switch vs. without	+0.0 [-1.3, +1.2]	-2.7 [-8.1, +2.5]	+0.0 [-1.1, +1.2]	-0.7 [-3.0, +1.5]	+0.0 [-1.5, +1.6]	+0.0 [-0.5, +0.4]
Laya							
	Peekaboo, head (1) vs. answer-key head (1)	+0.0 [-0.3, +0.3]	+0.0 [+0.0, +0.0]	+0.4 [+0.1, +0.7]	-4.8 [-7.5, -2.0]†	–	–
	Peekaboo, encoder (8) vs. their fine-tune	+1.6 [-2.1, +5.3]	+1.7 [-4.7, +8.8]	+10.9 [+4.7, +16.7]†	-7.6 [-9.8, -5.4]†	+21.6 [+19.4, +24.0]†	–
	Peekaboo, encoder (8) vs. answer-key encoder, same recipe	-2.1 [-5.1, +0.7]	+2.3 [-3.9, +7.9]	-1.4 [-5.1, +2.0]	-6.5 [-8.7, -4.1]†	-7.9 [-9.6, -6.1]†	–
	Peekaboo, encoder (8) vs. answer-key encoder, early-stopped	-1.8 [-4.7, +0.9]	-3.6 [-8.7, +1.4]	-0.9 [-4.5, +2.0]	-6.5 [-8.7, -4.3]†	-8.4 [-10.2, -6.7]†	–
	Peekaboo, encoder: 1 pass vs. 8 passes	-0.4 [-3.0, +2.1]	+1.4 [-4.4, +6.9]	-6.3 [-11.1, -1.9]†	-9.7 [-12.1, -7.1]†	-12.2 [-14.2, -10.3]†	–

Table 6: What the gain is made of. Left: gain over the released model of a per-task label prior with no input features (log-offsets on the released probabilities), trained from verdicts with Peekaboo's loop and from labels; and Peekaboo minus that prior. Right: how often eight on-policy passes reveal the gold answer: reward-1 events per training row (each one names the gold answer), the share of rows the base policy would reveal in 8 draws, the share of two-option rows, and the first pass's reward against the released model's best-answer accuracy on the same rows.

model	benchmark	prior (verdicts)	prior (labels)	Peekaboo − prior	params	reveals / row	P(revealed, 8 draws)	2-option rows	pass-1 reward	base best-answer acc.
Vega 0.8B	43 public tasks	+3.4	+2.3	+0.2 [-3.0, +3.1]	536	5.1	0.88	30%	0.50	0.55
	JevBench	–	–	–	–	4.7	–	–	0.49	–
	DecisionBench	+1.0	+2.7	+10.3 [+6.3, +14.7]	28584	4.9	0.76	21%	0.48	0.50
	typed-decisions	+4.3	+5.0	+1.9 [+0.1, +3.6]	71	6.1	0.94	30%	0.72	0.72
	emotion	+1.9	+3.1	+3.9 [+2.1, +5.7]	6	4.9	0.85	0%	0.57	0.56
	AG News	+1.4	+1.5	-0.3 [-1.2, +0.5]	4	7.5	0.97	0%	0.92	0.92
Vega 4B	43 public tasks	+0.1	-0.2	-1.2 [-3.6, +0.9]	536	6.1	0.92	30%	0.68	0.70
	JevBench	–	–	–	–	5.9	–	–	0.64	–
	DecisionBench	+1.0	+0.9	+4.7 [+2.0, +7.8]	28584	5.6	0.82	21%	0.62	0.61
	typed-decisions	+1.8	+3.9	-0.2 [-1.8, +1.6]	71	6.4	0.96	30%	0.75	0.79
	emotion	+2.5	+2.6	+2.9 [+1.1, +4.7]	6	4.8	0.84	0%	0.57	0.57
	AG News	+0.2	+0.2	+0.0 [-0.4, +0.6]	4	7.4	0.96	0%	0.92	0.93
Laya	typed-decisions	+15.3	+17.4	–	71	–	–	–	–	–
	emotion	-1.4	-3.3	–	6	–	–	–	–	–
	AG News	-0.6	-0.1	–	4	–	–	–	–	–

Table 7: Which rows change. Share of the released model's wrong rows that a method fixes, by the released model's rank of the gold option, and share of its right rows the method breaks (mean over benchmarks with a training split and over seeds).

model	method	fixed, all	gold 2nd	gold 3rd	gold 4th+	broken
Vega 0.8B	Peekaboo (8 passes)	32%	43%	23%	13%	11%
Vega 0.8B	Peekaboo (1 pass)	26%	37%	13%	9%	8%
Vega 0.8B	answer key (8 passes)	40%	48%	31%	23%	13%
Vega 4B	Peekaboo (8 passes)	25%	34%	15%	8%	8%
Vega 4B	Peekaboo (1 pass)	18%	27%	9%	4%	6%
Vega 4B	answer key (8 passes)	32%	40%	25%	16%	10%
Laya	Peekaboo (1 pass)	0%	1%	0%	0%	0%
Laya	Peekaboo, encoder (8)	47%	58%	45%	32%	16%
Laya	answer-key encoder (8)	55%	65%	55%	44%	15%

Table 8: VegaML's TTT head answers a row only where a head was kept for its task and the option set matches ("covered"); elsewhere its output is the released engine's. Eight-pass Peekaboo minus TTT, and the answer-key adapter minus TTT, on covered and uncovered rows, under the vegaml 0.12.3 rule used in the main comparison; rows covered under the 0.8.0 rule for reference.

model	benchmark	covered (0.12.3)	covered (0.8.0)	Peekaboo − TTT, covered	uncovered	answer key − TTT, covered
Vega 0.8B	43 public tasks	690 / 2535	1945	-1.9 [-8.6, +5.7]	+3.0 [-0.6, +6.9]	-2.7 [-8.3, +3.3]
Vega 0.8B	DecisionBench	329 / 1075	394	+3.2 [-2.0, +7.9]	+11.5 [+7.2, +16.1]	+3.9 [-1.3, +8.7]
Vega 0.8B	typed-decisions	600 / 2000	2000	-0.7 [-4.1, +2.8]	+1.7 [-0.2, +3.5]	-0.4 [-3.6, +2.9]
Vega 0.8B	emotion	1503 / 1503	1503	-6.4 [-8.4, -4.3]	–	-1.2 [-3.3, +0.9]
Vega 0.8B	AG News	0 / 1912	1912	–	+1.1 [+0.1, +2.0]	–
Vega 4B	43 public tasks	931 / 2535	2125	-3.7 [-8.5, +1.0]	-0.6 [-3.6, +2.4]	-2.8 [-6.9, +1.0]
Vega 4B	DecisionBench	302 / 1075	394	+3.1 [-1.9, +8.9]	+4.3 [+1.8, +7.1]	+3.6 [-1.9, +9.8]
Vega 4B	typed-decisions	600 / 2000	2000	+1.7 [-1.6, +5.1]	-2.0 [-4.4, +0.2]	+2.7 [-0.5, +6.0]
Vega 4B	emotion	1503 / 1503	1503	-6.8 [-8.9, -4.7]	–	-5.7 [-7.7, -3.7]
Vega 4B	AG News	1912 / 1912	1912	+1.5 [+0.4, +2.6]	–	+1.0 [-0.2, +2.1]

Table 9: Benchmark artifacts. Gain over the released model with every task, and without the five tasks with answer-position, option-structure or class-prior shortcuts (public: truthfulqa_mc1, pubmedqa; DecisionBench: contains_threat, toxicity_severity, finqa-numerical-reasoning).

model	method	public, all 43	public, 41 tasks	DecisionBench, all 43	DecisionBench, 40 tasks
Vega 0.8B	Peekaboo (8)	+3.7 [-0.1, +7.5]	+1.9 [-1.4, +5.1]	+11.3 [+7.0, +16.1]	+9.3 [+5.5, +14.0]
Vega 0.8B	Peekaboo (1)	+1.9 [-1.1, +5.2]	+0.5 [-2.1, +3.1]	+6.7 [+3.6, +10.5]	+5.9 [+2.8, +9.7]
Vega 0.8B	answer key	+3.4 [-0.9, +8.1]	+1.5 [-2.3, +5.3]	+12.3 [+7.8, +17.5]	+10.2 [+6.0, +15.1]
Vega 0.8B	TTT (0.12.3)	+2.0 [-0.8, +5.6]	+1.0 [-1.4, +3.5]	+2.3 [-1.3, +6.8]	+1.3 [-1.9, +5.5]
Vega 4B	Peekaboo (8)	-1.1 [-3.8, +1.5]	-1.8 [-4.5, +0.8]	+5.7 [+2.7, +9.3]	+4.6 [+2.1, +7.1]
Vega 4B	Peekaboo (1)	+0.0 [-2.9, +2.6]	-0.4 [-3.6, +2.1]	+3.2 [+1.5, +5.4]	+3.4 [+1.6, +5.7]
Vega 4B	answer key	-1.2 [-4.1, +1.8]	-2.2 [-5.0, +0.6]	+5.1 [+1.4, +9.1]	+3.7 [+0.5, +7.3]
Vega 4B	TTT (0.12.3)	+0.6 [-1.6, +3.0]	+0.1 [-2.1, +2.2]	+1.8 [-0.2, +4.2]	+1.3 [-0.6, +3.5]
Laya	Peekaboo (1)	+0.1 [-0.1, +0.3]	+0.1 [-0.1, +0.3]	+0.1 [+0.0, +0.3]	+0.1 [+0.0, +0.3]
Laya	Peekaboo, encoder	+2.6 [-0.8, +6.0]	+2.0 [-1.4, +5.5]	+14.0 [+8.8, +19.2]	+12.9 [+7.7, +18.1]
Laya	their fine-tune	+1.0 [-1.2, +3.5]	+0.8 [-1.6, +3.4]	+3.1 [-0.4, +7.3]	+2.7 [-1.1, +7.1]

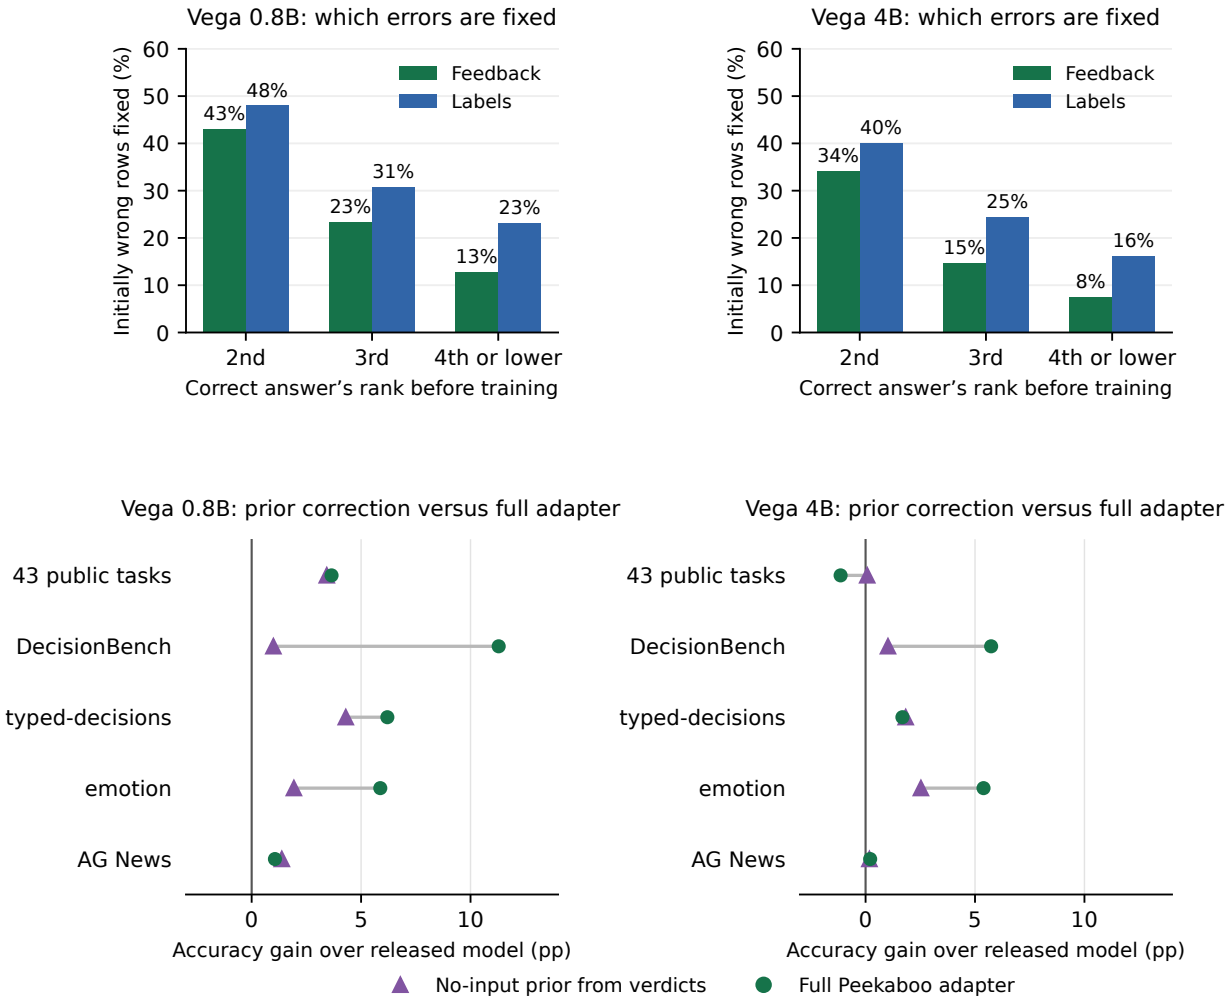


Figure 6: Two explanations of the gain, with model sizes in separate columns. Top: percentage of initially wrong rows corrected after eight passes, by the correct answer's initial rank; bars distinguish feedback and same-recipe labelled training. Values are unweighted means across non-JevBench benchmarks with a defined rate. Bottom: a no-input prior learned from verdicts (purple triangles) versus the full Peekaboo adapter (green circles), as gains over the released model. Connecting lines compare point estimates, not uncertainty or training trajectories. Complete prior controls and uncertainty appear in Table 6.

Peekaboo's +6.2 and +1.7; Peekaboo minus the prior is +1.9 [+0.1, +3.6] and −0.2 [−1.8, +1.6]. About two thirds of the typed-decisions gain at 0.8B, and all of it at 4B, is a prior shift that 71 numbers can learn from verdicts. On the public menu the prior gains +3.4 and +0.1 (Peekaboo minus prior +0.2, −1.2). On DecisionBench the prior gains +1.0 and +1.0 against Peekaboo's +11.3 and +5.7, so there the adapter learns something the option text alone does not give (Peekaboo minus prior +10.3 [+6.3, +14.7], +4.7 [+2.0, +7.8]); on emotion the prior gains +1.9 and +2.5 against +5.9 and +5.4. DecisionBench and emotion are the benchmarks on which the paper's claim about learning from verdicts rests; typed-decisions and the public menu are largely prior correction, which verdicts also deliver, but which the per-task offset model can also deliver (71 parameters on typed-decisions and 536 on the public menu).

Which rows change. Table 7 and Figure 6 (top) test the prediction of the $\pi(y)$ weighting. Over the benchmarks with a training split, eight-pass Peekaboo fixes 32% of the released model's wrong rows at 0.8B and

Table 10: Calibration, proper scores and selective risk (lower is better; mean over seeds). "+ T" rows apply one temperature to the released model, fitted on its own calibration rows from verdicts on sampled answers (mean over 10 draws) or from labels; JevBench has no calibration pool for them.

model	method	public				JevBench				DecisionB.				typed-dec.				emotion				AG News			
		ECE	NLL	Brier	AURC	ECE	NLL	Brier	AURC	ECE	NLL	Brier	AURC	ECE	NLL	Brier	AURC	ECE	NLL	Brier	AURC	ECE	NLL	Brier	AURC
Vega 0.8B																									
	released model	0.150	1.100	0.550	0.363	0.184	1.262	0.619	0.281	0.101	1.631	0.638	0.291	0.063	0.753	0.446	0.200	0.031	1.193	0.575	0.274	0.021	0.291	0.149	0.023
	released + T from verdicts	0.151	1.090	0.549	0.363	–	–	–	–	0.072	1.509	0.632	0.294	0.053	0.747	0.444	0.200	0.033	1.192	0.575	0.274	0.016	0.289	0.149	0.023
	released + T from labels	0.148	1.088	0.548	0.363	–	–	–	–	0.089	1.515	0.637	0.296	0.050	0.745	0.443	0.201	0.036	1.191	0.575	0.274	0.016	0.289	0.148	0.023
	Peekaboo before the honesty fix	0.308	2.735	0.670	0.321	0.275	1.758	0.709	0.308	0.257	3.547	0.620	0.186	0.200	1.354	0.462	0.151	0.293	3.593	0.646	0.215	0.073	0.707	0.159	0.026
	Peekaboo	0.187	1.114	0.548	0.322	0.093	1.036	0.576	0.312	0.053	1.316	0.522	0.185	0.055	0.690	0.387	0.150	0.052	1.057	0.512	0.213	0.023	0.281	0.139	0.027
	Peekaboo, one pass	0.177	1.102	0.553	0.342	0.079	1.006	0.560	0.287	0.059	1.376	0.553	0.212	0.059	0.714	0.405	0.160	0.046	1.045	0.511	0.211	0.012	0.256	0.132	0.020
	answer key + label temperature	0.174	1.073	0.536	0.323	0.086	1.053	0.572	0.304	0.075	1.178	0.493	0.155	0.076	0.685	0.376	0.135	0.033	0.911	0.444	0.151	0.024	0.291	0.151	0.023
Vega 4B																									
	released model	0.146	0.768	0.383	0.169	0.162	0.959	0.459	0.166	0.095	1.423	0.496	0.157	0.051	0.631	0.367	0.126	0.056	1.167	0.565	0.251	0.025	0.224	0.113	0.012
	released + T from verdicts	0.147	0.761	0.387	0.169	–	–	–	–	0.053	1.206	0.484	0.158	0.066	0.643	0.370	0.126	0.019	1.148	0.561	0.250	0.017	0.214	0.112	0.012
	released + T from labels	0.153	0.764	0.389	0.169	–	–	–	–	0.067	1.203	0.488	0.159	0.053	0.632	0.367	0.126	0.019	1.148	0.560	0.250	0.012	0.211	0.111	0.012
	Peekaboo before the honesty fix	0.244	2.374	0.505	0.180	0.192	1.316	0.476	0.166	0.237	3.328	0.529	0.132	0.185	1.237	0.431	0.121	0.283	2.981	0.640	0.232	0.046	0.407	0.119	0.012
	Peekaboo	0.167	0.836	0.412	0.179	0.076	0.783	0.424	0.174	0.047	1.095	0.432	0.129	0.051	0.637	0.359	0.118	0.049	1.087	0.520	0.226	0.012	0.208	0.107	0.012
	Peekaboo, one pass	0.143	0.775	0.386	0.171	0.064	0.776	0.433	0.178	0.040	1.135	0.455	0.144	0.055	0.643	0.360	0.123	0.028	1.115	0.535	0.235	0.013	0.208	0.108	0.013
	answer key + label temperature	0.173	0.840	0.421	0.191	0.096	0.679	0.383	0.131	0.062	1.051	0.429	0.116	0.056	0.567	0.324	0.100	0.054	0.996	0.491	0.187	0.011	0.220	0.117	0.013
Laya																									
	released model	0.217	1.651	0.599	0.365	0.125	0.928	0.534	0.236	0.165	2.062	0.748	0.454	0.182	1.361	0.770	0.610	0.303	2.082	0.685	0.267	0.015	0.206	0.107	0.012
	released + T from verdicts	0.171	1.189	0.561	0.365	–	–	–	–	0.110	1.631	0.707	0.404	0.058	1.207	0.682	0.590	0.095	1.176	0.566	0.282	0.016	0.206	0.107	0.012
	released + T from labels	0.174	1.159	0.574	0.361	–	–	–	–	0.103	1.631	0.705	0.404	0.059	1.205	0.684	0.588	0.097	1.177	0.566	0.282	0.017	0.206	0.108	0.012
	Peekaboo, one pass	0.163	1.111	0.553	0.364	0.107	0.911	0.521	0.245	0.075	1.575	0.692	0.389	0.034	1.194	0.675	0.563	–	–	–	–	–	–	–	–
	Peekaboo, encoder unfrozen	0.157	1.088	0.536	0.329	0.115	1.024	0.570	0.309	0.052	1.392	0.574	0.223	0.077	0.827	0.460	0.213	0.094	0.665	0.289	0.070	–	–	–	–
	answer-key encoder + label temperature	0.179	1.025	0.520	0.303	0.148	1.085	0.612	0.448	0.057	1.296	0.543	0.211	0.045	0.636	0.360	0.129	0.035	0.313	0.153	0.018	–	–	–	–

breaks 11% of its right ones; the labelled adapter fixes 40% and breaks 13%. Where the released model ranked the gold option second, Peekaboo fixes 43% of rows against the adapter's 48%; where it ranked it fourth or lower, 13% against 23%. The verdict-trained adapter is the more conservative learner: it fixes fewer of the deep errors, consistent with the $\pi(y)$-weighted expected gradient, and breaks fewer correct answers. The two differences partly offset each other, which is why the net accuracies are close; at DecisionBench's base accuracy the fixing deficit is the larger of the two.

Artifacts. Table 9 repeats the public-menu and DecisionBench gains without the five tasks with answer-position, option-structure or class-prior shortcuts (two on the public menu, three on DecisionBench). On the public menu at 0.8B Peekaboo's gain goes from +3.7 to +1.9 $[-1.4, +5.1]$; at 4B from −1.1 to −1.8. On DecisionBench from +11.3 to +9.3 $[+5.5, +14.0]$ and from +5.7 to +4.6 $[+2.1, +7.1]$. The labelled adapter and TTT profit from the same tasks, so the comparisons between methods move little; the gains over the released model are what the artifacts inflate. The per-task tables in Appendix D show where the rest of the DecisionBench gain sits.

6.4 Calibration

Before the honesty fix the trained add-on is badly over-confident: ECE 0.257 on DecisionBench at 0.8B, 0.237 at 4B, 0.293 and 0.283 on emotion (Table 10, Figure 8). One temperature fitted from verdicts brings these to 0.053, 0.047, 0.052 and 0.049. Against the released model the picture is mixed. Peekaboo's ECE is lower on DecisionBench and JevBench at both sizes, on typed-decisions at 0.8B and on emotion and AG News at 4B, and higher on the public menu at both sizes and on emotion and AG News at 0.8B (Table 10). The released model with the same temperature fitted on its own calibration rows reaches 0.072 and 0.053 on DecisionBench against the released model's 0.101 and 0.095; Peekaboo is better calibrated than this control in four of the ten Vega cells where both exist (DecisionBench at both sizes, typed-decisions and AG News at 4B) and worse in the other six. This control shows that temperature fitting accounts for much of the ECE improvement relative to the unscaled released model. On proper scores the adapter does help: its Brier score is below the released model's on every Vega benchmark but the public menu at 4B, and on DecisionBench at 0.8B its NLL is 1.316 against the released model's 1.631, though above the labelled adapter's 1.178. Selective risk (AURC), which depends almost only on how rows are ranked by confidence, is better than the released model's on DecisionBench, typed-decisions and emotion at both sizes and on the public menu at



Figure 7: Confidence quality with values printed directly. Columns distinguish the released base model, base with a verdict-fitted temperature, Peekaboo before and after temperature fitting (Pk raw and Pk + T), and the labelled adapter with a label-fitted temperature. ECE is left and selective risk (AURC) right; lower is better. Darker shading indicates larger error with a common scale across models within each metric. Gray cells are unavailable, not zeros. Values are rounded means; Table 10 gives more precision and proper scores. Laya uses encoder training.

0.8B, and worse on JevBench at both sizes, on the public menu at 4B and on AG News at 0.8B (Figure 7). The fitted temperatures range from 1.3 to 10.3 with a median of 3.8 (Figure 9); part of each is inherited from the released model, which itself prefers a temperature above 1 on most benchmarks (2.0 on DecisionBench at 0.8B).

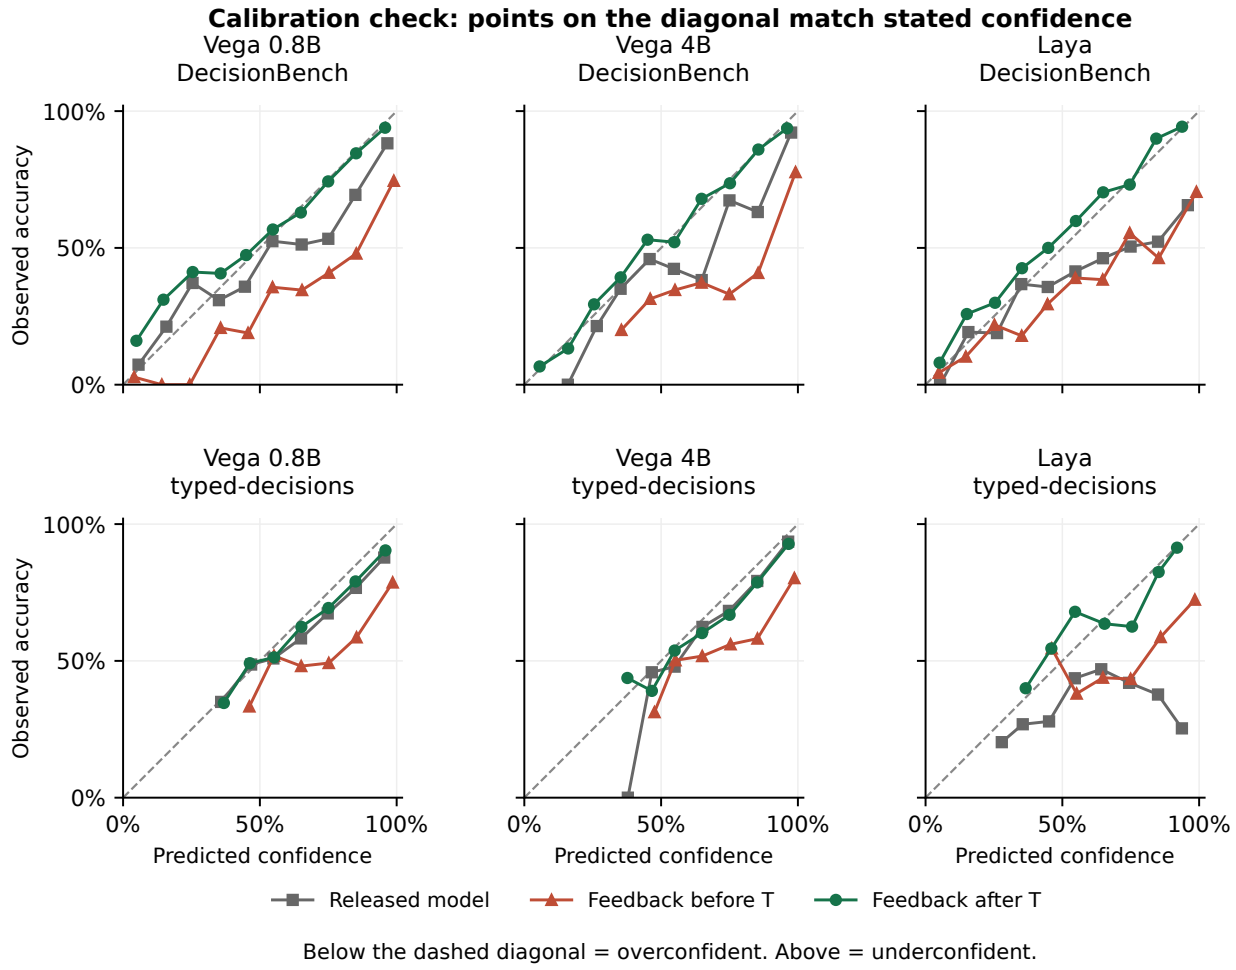


Figure 8: How to read confidence. Each point groups predictions in one of ten confidence bins, plotting mean confidence against observed accuracy. The dashed diagonal is perfect calibration; below means overconfidence. Colors and marker shapes distinguish released, feedback-trained and temperature-scaled policies. Test predictions are pooled across available seeds, with at least ten predictions per displayed bin; unanswered rows without probabilities are excluded. These ten-bin diagrams are descriptive and differ from the fifteen-bin ECE in Table 10. Laya uses encoder training.

6.5 The KL term and the switch

Removing the KL term changes eight-pass Peekaboo by +1.2 [−0.6, +3.0] points on DecisionBench at 0.8B and +2.2 [+0.4, +4.2] at 4B (positive means the term helps; Table 5), +1.9 and +0.4 on typed-decisions, +1.1 and −0.6 on the public menu, and +0.5 and +11.0 [+8.5, +13.5] on emotion. The term is decisive in one cell, emotion at 4B, where the run without it collapsed, and that is the only KL contrast that survives Holm adjustment; elsewhere it is worth between −0.6 and +2.2 points and none of those differences survives. The no-KL arm has three seeds on DecisionBench and on JevBench at 0.8B, and one elsewhere. The single-seed emotion result does not establish that collapse is typical without KL regularisation.

The switch compares best answers, which takes two verdicts on the same question, one of them on an answer the system did not serve; a deployment would have to run it as a held-out comparison on a slice of traffic. At 0.8B it turns the add-on off on JevBench, where Peekaboo loses 2.3 points (gain with the switch +0.0), and keeps it on every other benchmark. At 4B it keeps the add-on on DecisionBench and

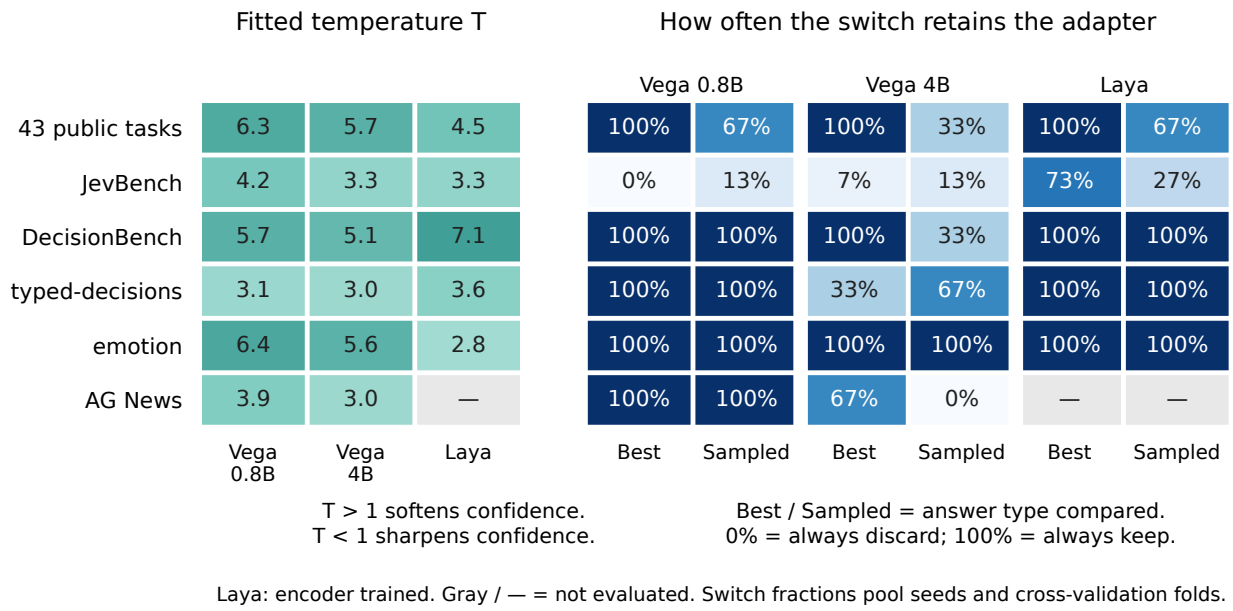


Figure 9: Two post-training decisions. Left: mean fitted temperature, printed as T; values above one soften probabilities and below one sharpen them (shading uses $\log_2 T$ centered at one). Right: percentage of calibration comparisons retaining the adapter, with best-answer and sampled-answer rules separated for each model. Zero means always discarded and 100% always retained in the observed comparisons. Fractions pool training seeds and, on JevBench, cross-validation folds; they are not independent deployment trials. Gray cells were not evaluated. Laya uses encoder training.

emotion; it also keeps it on the public menu, where Peekaboo is −1.1 (not significant), and it drops it in most JevBench folds and on typed-decisions, giving up gains of +2.9 and +1.7 that were themselves not significant. The sampled-answer switch diagnostic, computed on the same runs, kept the add-on in 33% of the DecisionBench runs at 4B, where the best-answer switch keeps it in 100%; the two statistics disagree in the direction Section 4.5 describes.

6.6 Laya: head, encoder, and the matched fine-tune

With the encoder frozen, a trainable copy of Laya's head learns a prior shift and little else; in a historical grid, eight passes gave +13.2 points on typed-decisions where a 71-offset prior from the same verdicts gave +15.4, and in the current grid one pass gives +0.2. With the encoder unfrozen, verdicts move the whole model: +32.3 [+29.0, +35.5] on typed-decisions with eight passes, +22.6 with one, +22.2 on emotion, +14.0 on DecisionBench and +2.6 on the public menu. The matched labelled comparison is the same fine-tune trained by cross-entropy, same rows, passes and KL term, encoder unfrozen. It gains +38.8 on typed-decisions, +30.0 on emotion, +15.4 on DecisionBench and +4.6 on the public menu; Peekaboo minus this arm is −6.5 [−8.7, −4.1] on typed-decisions, −7.9 [−9.6, −6.1] on emotion, −1.4 on DecisionBench and −2.1 on the public menu. At matched capacity and data, labels outperform eight rounds of verdicts on typed-decisions and emotion. These encoder comparisons have one seed each, so their intervals describe test-row uncertainty conditional on those runs. On DecisionBench and the public menu the point-estimate gaps are one to two points and the intervals include zero; this does not establish equal performance. On JevBench the differences are not significant: Peekaboo minus the same-recipe arm is +2.3 [−3.9, +7.9] and minus the early-stopped one −3.6 [−8.7, +1.4]. One pass costs the encoder arm a lot (−9.7 on typed-decisions, −12.2 on emotion against eight passes); no one-pass labelled encoder arm was run, so these have no matched label comparator.

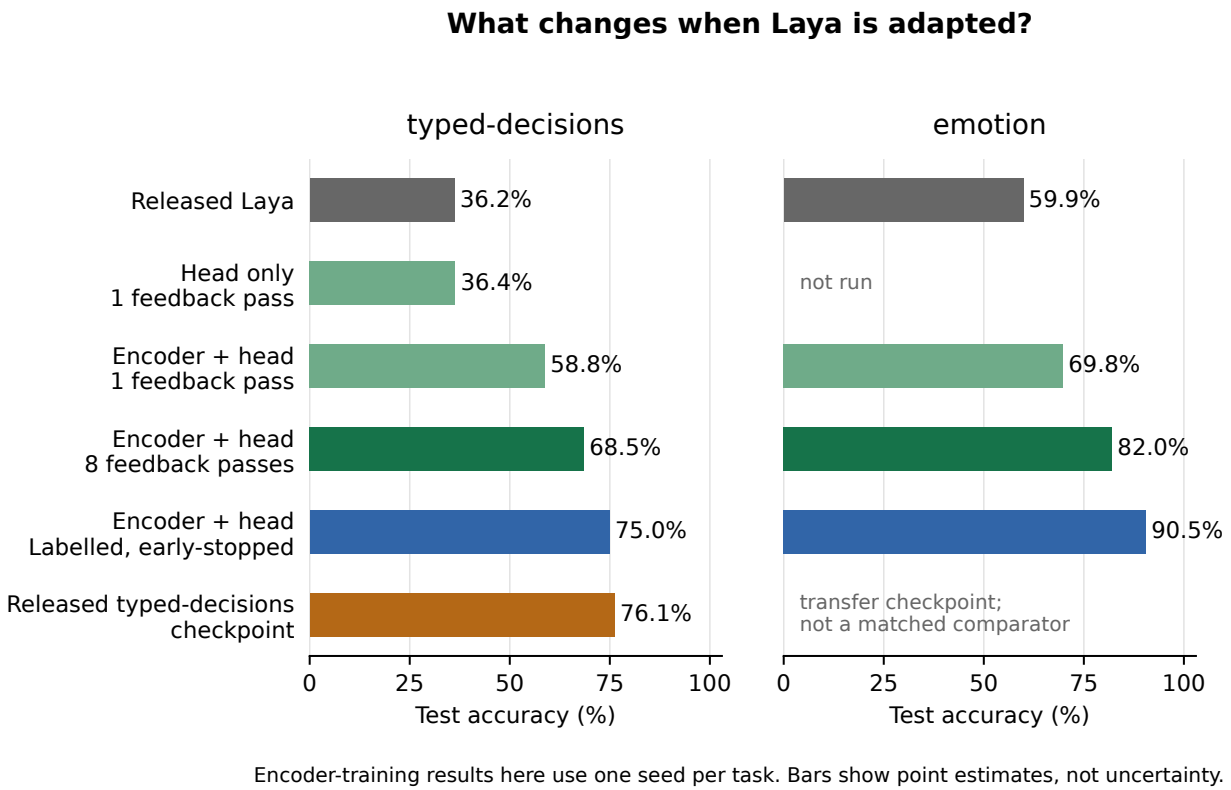


Figure 10: Laya adaptation on typed-decisions and emotion, shown as test accuracy. Named rows distinguish released Laya, head-only feedback, encoder feedback and the matched labelled encoder with early stopping. Encoder results here each use one seed; bars show point estimates, not confidence intervals. The released typed-decisions checkpoint is included only in-domain; its emotion transfer result is not a matched comparator. Head-only emotion training was not run. All six benchmarks and paired intervals remain in Tables 2 and 5.

The head-only one-pass arm barely moves from the released model (+0.2 on typed-decisions) where the one-pass labelled head gains +5.0. On typed-decisions the matched fine-tune reaches 0.750 against the released `laya-typed-decisions` checkpoint's 0.761, so our labelled arm is a fair stand-in for Laya's own recipe. Away from typed-decisions the released fine-tune is a transfer model that never saw the benchmark, and in-domain training of any kind beats it: Peekaboo with the encoder is +21.6 above it on emotion and +10.9 on DecisionBench, the labelled encoder arm +29.5 and +12.3. Those numbers say that fine-tuning on a benchmark's own rows helps on that benchmark; they say nothing about verdicts against labels, and we do not use them to.

6.7 typed-decisions without Vega's shipped adapter

Vega's released `typed_decisions` adapter was fitted on the pool our methods learn from, so on this benchmark "released Vega" already carries those labels. With the shipped adapters switched off the released engine scores 0.377 at 0.8B and 0.440 at 4B (Table 11). From that start eight-pass Peekaboo gains +13.0 [+9.8, +16.5] and +24.6 [+21.0, +28.5], the labelled adapter +31.6 and +30.8, and one-pass Peekaboo +12.1 and +22.2. The gap between verdicts and labels, −18.6 at 0.8B and −6.3 at 4B, is far wider than with the shipped adapter on, which is the clearest demonstration in the paper of what the $\pi(y)$ weighting means: when the released model puts little probability on the gold answer, a verdict on a sampled answer rarely finds it, and labels do.

Table 11: typed-decisions on Vega with the shipped `typed_decisions` adapter, which was fitted with labels on this feedback pool, switched off. Accuracy (mean over seeds) and gain over the corresponding released model.

model	method	accuracy, adapters on	gain	accuracy, adapters off	gain
Vega 0.8B	released engine	0.666	–	0.377	–
Vega 0.8B	answer-key adapter (8)	0.737 ±0.002	+7.0 [+5.0, +9.1]	0.693 ±0.013	+31.6 [+28.9, +34.4]
Vega 0.8B	Peekaboo (8)	0.728 ±0.007	+6.2 [+4.2, +8.2]	0.507 ±0.028	+13.0 [+9.8, +16.5]
Vega 0.8B	Peekaboo (1)	0.708 ±0.003	+4.1 [+2.2, +6.1]	0.498 ±0.007	+12.1 [+10.1, +14.2]
Vega 4B	released engine	0.729	–	0.440	–
Vega 4B	answer-key adapter (8)	0.769 ±0.001	+4.0 [+2.5, +5.6]	0.748 ±0.006	+30.8 [+27.8, +33.9]
Vega 4B	Peekaboo (8)	0.746 ±0.010	+1.7 [-0.1, +3.6]	0.685 ±0.027	+24.6 [+21.0, +28.5]
Vega 4B	Peekaboo (1)	0.749 ±0.004	+2.0 [+0.7, +3.5]	0.661 ±0.004	+22.2 [+19.4, +25.0]

7 Limitations

- **Simulated correctness and collection costs.** Verdicts come from benchmark gold answers. Flipped, missing and complaint-skewed verdicts are simulations, not measurements of human behaviour. Real feedback studies [15, 29] motivate testing systematic rater errors and selection bias. Noise is tested only under repeated feedback, not the one-rating setting. Additional calibration and switch verdicts mean the pipeline is not a one-verdict total annotation budget. We do not measure live exploration cost, user outcomes or annotation time.
- **Seeds and incomplete comparisons.** Table 12 lists the available independent training runs. The main Vega feedback and labelled arms have three seeds per benchmark. Laya encoder arms have three on the public menu and JevBench, two on DecisionBench and one on typed-decisions and emotion; AG News and a matched one-pass labelled encoder arm were not run. The realistic-feedback arms have three seeds on DecisionBench and one on typed-decisions and emotion. No-KL arms have three on DecisionBench and Vega 0.8B JevBench, and one elsewhere. Head-only Laya one-pass arms have three on typed-decisions and one on the public menu, JevBench and DecisionBench. Historical head-only eight-pass results are identified separately; their labelled head lacked the KL term. Single-seed intervals cannot estimate variation across training runs. The failed half-rated emotion run at 4B is retained and disclosed in Section 6.1.
- **Model and modality scope.** The study covers two released model families from one developer. Every evaluated task uses the text input path, including structured states rendered as text. Vega's image capability is not evaluated. A frozen vision encoder does not guarantee preserved image decisions when a shared downstream adapter changes; neither multimodal improvement nor retention is established. JevBench and DecisionBench are small, and the released models are already near 0.93 on AG News.
- **Benchmark artifacts and overlap.** Five tasks have answer-position, option-structure or class-prior shortcuts. Vega's shipped typed-decisions adapter was fitted on the feedback pool; the engines' original training corpora are undisclosed. Removing the identified tasks and disabling shipped adapters are sensitivity analyses, not proof that every remaining task is free of contamination.
- **Historical design choices.** The KL coefficient was selected in a pilot involving DecisionBench test rows (Appendix A), so that benchmark is not an untouched test of the full research process. Rank, learning rates and pass count were fixed before this grid. Labelled arms use calibration labels for early stopping; the main feedback arms do not receive that model-selection signal. The grid was reduced while running to fit the GPU budget (Appendix B).
- **Calibration and statistical precision.** One temperature per benchmark may be too coarse for a heterogeneous task menu. ECE is sensitive to binning and sample size, so proper scores are also reported. Non-significant differences are inconclusive rather than evidence of equality; the exploratory non-inferiority marks use unadjusted intervals. Three seeds give limited information about training variability, and single-seed intervals are conditional on the observed run.

- **Reproduction and causal attribution.** Historical same-seed reruns differ across machines by up to about 3 points. The complete current grid has not been independently reproduced. Comparing logged and fresh on-policy feedback changes both information and the learning procedure; it does not identify a pure causal effect of verdict count.

8 Conclusion

Correctness feedback can improve the evaluated small text decision models, but its value depends on the feedback budget, the initial policy and the adaptation method. On DecisionBench at 0.8B, one on-policy training pass gains +6.7 points, reusing one logged sampled verdict per row for eight optimisation passes gains +8.3, and eight passes with fresh verdicts gain +11.3. The matched early-stopped labelled adapter gains +13.2. Repeated feedback often reveals labels, and the expected reward gradient is the supervised gradient weighted by the current probability of the correct answer. The logged comparison is consistent with a benefit from additional verdicts, but changes in policy and objective prevent a purely causal attribution to feedback count.

The controls narrow what the gains mean. Prior correction explains much of the improvement on typed-decisions and the public menu; matched labelled fine-tuning outperforms feedback on Laya where the differences are clearest; and a temperature-fitted released model explains much of the calibration gain. Benchmark artifacts, pre-existing training-pool overlap, limited seeds and historical hyper-parameter selection involving test rows restrict generalisation. This study establishes results for simulated feedback on the evaluated text tasks, not for human feedback, image inputs or unrestricted deployments. The next evidence needed is an independent reproduction of the main comparisons and evaluation with one naturally collected verdict per served answer, followed by modality-specific tests if multimodal claims are intended.

Acknowledgements. We thank the authors of VegaML and Laya for releasing their models, benchmark code and training recipes.

References

[1] Arash Ahmadian, Chris Cremer, Matthias Gallé, Marzieh Fadaee, Julia Kreutzer, Olivier Pietquin, Ahmet Üstün, and Sara Hooker. Back to basics: revisiting REINFORCE-style optimization for learning from human feedback in LLMs. In *Proceedings of ACL*, 2024.

[2] Nishant Balepur, Abhilasha Ravichander, and Rachel Rudinger. Artifacts or abduction: how do LLMs answer multiple-choice questions without the question? In *Proceedings of ACL*, 2024.

[3] Jon Ander Campos, Kyunghyun Cho, Arantxa Otegi, Aitor Soroa, Gorka Azkune, and Eneko Agirre. Improving conversational question answering systems after deployment using feedback-weighted learning. In *Proceedings of COLING*, 2020.

[4] Leshem Choshen, Lior Fox, Zohar Aizenbud, and Omri Abend. On the weaknesses of reinforcement learning for neural machine translation. In *Proceedings of ICLR*, 2020.

[5] Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*, 2017.

[6] Convai Innovations. Laya: fast, non-autoregressive System 1 decision engine with calibrated probabilities. https://huggingface.co/convaiinnovations/laya, 2026. Python package `laya` 0.3.23; checkpoints `convaiinnovations/laya` and `convaiinnovations/laya-typed-decisions`.

[7] Koby Crammer and Claudio Gentile. Multiclass classification with bandit feedback using adaptive regularization. In *Proceedings of ICML*, 2011.

[8] Mehul Damani, Isha Puri, Stewart Slocum, Idan Shenfeld, Leshem Choshen, Yoon Kim, and Jacob Andreas. Beyond binary rewards: training LMs to reason about their uncertainty. In *Proceedings of ICLR*, 2026.

[9] Amit Daniely and Tom Helbertal. The price of bandit information in multiclass online classification. In *Proceedings of COLT*, 2013.

[10] Damek Davis and Benjamin Recht. Reinforcement learning with binary rewards is gradient ascent on a transform of the success probability. *arXiv preprint arXiv:2510.13651*, 2025.

[11] Hanze Dong, Wei Xiong, Deepanshu Goyal, Yihan Zhang, Winnie Chow, Rui Pan, Shizhe Diao, Jipeng Zhang, Kashun Shum, and Tong Zhang. RAFT: reward ranked finetuning for generative foundation model alignment. *Transactions on Machine Learning Research*, 2023.

[12] Miroslav Dudík, John Langford, and Lihong Li. Doubly robust policy evaluation and learning. In *Proceedings of ICML*, 2011.

[13] Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. KTO: model alignment as prospect theoretic optimization. In *Proceedings of ICML*, 2024.

[14] Ge Gao, Eunsol Choi, and Yoav Artzi. Simulating bandit learning from user feedback for extractive question answering. In *Proceedings of ACL*, 2022.

[15] Ge Gao, Hung-Ting Chen, Yoav Artzi, and Eunsol Choi. Continually improving extractive QA via human feedback. In *Proceedings of EMNLP*, 2023.

[16] Aritra Ghosh, Himanshu Kumar, and P. S. Sastry. Robust loss functions under label noise for deep neural networks. In *Proceedings of AAAI*, 2017.

[17] Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.

[18] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of ICML*, 2017.

[19] Chirag Gupta and Aaditya Ramdas. Top-label calibration and multiclass-to-binary reductions. In *Proceedings of ICLR*, 2022.

[20] Braden Hancock, Antoine Bordes, Pierre-Emmanuel Mazaré, and Jason Weston. Learning from dialogue after deployment: feed yourself, chatbot! In *Proceedings of ACL*, 2019.

[21] Hanno Labs. DecisionBench. https://huggingface.co/datasets/Hanno-Labs/decision-bench, 2026.

[22] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *Proceedings of ICLR*, 2022.

[23] Natasha Jaques, Asma Ghandeharioun, Judy Hanwen Shen, Craig Ferguson, Agata Lapedriza, Noah Jones, Shixiang Gu, and Rosalind Picard. Way off-policy batch deep reinforcement learning of implicit human preferences in dialog. In *arXiv preprint arXiv:1907.00456*, 2019.

[24] Thorsten Joachims, Adith Swaminathan, and Maarten de Rijke. Deep learning with logged bandit feedback. In *Proceedings of ICLR*, 2018.

[25] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Dodds, Andy Jones, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

[26] Sham M. Kakade, Shai Shalev-Shwartz, and Ambuj Tewari. Efficient bandit algorithms for online multiclass prediction. In *Proceedings of ICML*, 2008.

[27] Samuel Kiegeland and Julia Kreutzer. Revisiting the weaknesses of reinforcement learning for neural machine translation. In *Proceedings of NAACL*, 2021.

[28] Julia Kreutzer, Artem Sokolov, and Stefan Riezler. Bandit structured prediction for neural sequence-to-sequence learning. In *Proceedings of ACL*, 2017.

[29] Julia Kreutzer, Shahram Khadivi, Evgeny Matusov, and Stefan Riezler. Can neural machine translation be improved with user feedback? In *Proceedings of NAACL (Industry Track)*, 2018.

[30] Julia Kreutzer, Joshua Uyheng, and Stefan Riezler. Reliability and learnability of human bandit feedback for sequence-to-sequence reinforcement learning. In *Proceedings of ACL*, 2018.

[31] Ananya Kumar, Percy Liang, and Tengyu Ma. Verified uncertainty calibration. In *Advances in Neural Information Processing Systems*, 2019.

[32] Ananya Kumar, Aditi Raghunathan, Robbie Jones, Tengyu Ma, and Percy Liang. Fine-tuning can distort pretrained features and underperform out-of-distribution. In *Proceedings of ICLR*, 2022.

[33] Carolin Lawrence and Stefan Riezler. Improving a neural semantic parser by counterfactual learning from human bandit feedback. In *Proceedings of ACL*, 2018.

[34] Zachary C. Lipton, Yu-Xiang Wang, and Alexander J. Smola. Detecting and correcting for label shift with black box predictors. In *Proceedings of ICML*, 2018.

[35] LocalLLaMA. typed-decisions. `https://huggingface.co/datasets/LocalLLaMA/typed-decisions`, 2026.

[36] Mahdi Pakdaman Naeini, Gregory F. Cooper, and Milos Hauskrecht. Obtaining well calibrated probabilities using Bayesian binning into quantiles. In *Proceedings of AAAI*, 2015.

[37] NandhaKishorM. VegaML: typed decisions from a frozen language model and a small physics engine. `https://github.com/NandhaKishorM/vegaml`, 2026. Python package `vegaml` 0.8.0; models `nandakishorm/vega-08b-public-intents` and the 4B release.

[38] Khanh Nguyen, Hal Daumé III, and Jordan Boyd-Graber. Reinforcement learning for bandit neural machine translation with simulated human feedback. In *Proceedings of EMNLP*, 2017.

[39] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, 2022.

[40] Qwen Team. Qwen3.5-4B model card. https://huggingface.co/Qwen/Qwen3.5-4B, 2026.

[41] Qwen Team. Qwen3.5-0.8B model card. https://huggingface.co/Qwen/Qwen3.5-0.8B, 2026.

[42] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of ICML*, 2021.

[43] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, 2023.

[44] Rebecca Roelofs, Nicholas Cain, Jonathon Shlens, and Michael C. Mozer. Mitigating bias in calibration error estimation. In *Proceedings of AISTATS*, 2022.

[45] Marco Saerens, Patrice Latinne, and Christine Decaestecker. Adjusting the outputs of a classifier to new a priori probabilities: a simple procedure. *Neural Computation*, 14(1):21–41, 2002.

[46] Elvis Saravia, Hsien-Chi Toby Liu, Yen-Hao Huang, Junlin Wu, and Yi-Shin Chen. CARER: Contextualized affect representations for emotion recognition. In *Proceedings of EMNLP*, 2018.

[47] Idan Shenfeld, Jyothish Pari, and Pulkit Agrawal. RL's razor: why online reinforcement learning forgets less. In *Proceedings of ICLR*, 2026.

[48] Artem Sokolov, Julia Kreutzer, Christopher Lo, and Stefan Riezler. Learning structured predictors from bandit feedback for interactive NLP. In *Proceedings of ACL*, 2016.

[49] Florian Standhartinger. JevBench: public tiers. https://huggingface.co/datasets/fstandhartinger/jevbench, 2026.

[50] Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel M. Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F. Christiano. Learning to summarize from human feedback. In *Advances in Neural Information Processing Systems*, 2020.

[51] Adith Swaminathan and Thorsten Joachims. Counterfactual risk minimization: learning from logged bandit feedback. In *Proceedings of ICML*, 2015.

[52] Yuanzhi Wang and Xinyu Qiao. Bandit-feedback online multiclass classification with conformal prediction. In *Proceedings of ICML*, 2024.

[53] Benjamin Warner, Antoine Chaffin, Benjamin Clavié, Orion Weller, Oskar Hallström, Said Taghadouini, et al. Smarter, better, faster, longer: a modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference. *arXiv preprint arXiv:2412.13663*, 2024.

[54] Mitchell Wortsman, Gabriel Ilharco, Jong Wook Kim, Mike Li, Simon Kornblith, Rebecca Roelofs, Raphael Gontijo Lopes, Hannaneh Hajishirzi, Ali Farhadi, Hongseok Namkoong, and Ludwig Schmidt. Robust fine-tuning of zero-shot models. In *Proceedings of CVPR*, 2022.

[55] Wei Xiong, Jiarui Yao, Yuhui Xu, Bo Pang, Lei Wang, Doyen Sahoo, Junnan Li, Nan Jiang, Tong Zhang, Caiming Xiong, and Hanze Dong. A minimalist approach to LLM reasoning: from rejection sampling to REINFORCE. *arXiv preprint arXiv:2504.11343*, 2025.

[56] Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang. Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model? In *Advances in Neural Information Processing Systems*, 2025.

[57] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems*, 2022.

[58] Xiang Zhang, Junbo Zhao, and Yann LeCun. Character-level convolutional networks for text classification. In *Advances in Neural Information Processing Systems*, 2015.

[59] Zhilu Zhang and Mert R. Sabuncu. Generalized cross entropy loss for training deep neural networks with noisy labels. In *Advances in Neural Information Processing Systems*, 2018.

[60] Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. Calibrate before use: improving few-shot performance of language models. In *Proceedings of ICML*, 2021.

[61] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *Proceedings of ICLR*, 2024.

[62] Han Zhou, Xingchen Wan, Lev Proleev, Diana Mincu, Jilin Chen, Katherine Heller, and Subhrajit Roy. Batch calibration: rethinking calibration for in-context learning and prompt engineering. In *Proceedings of ICLR*, 2024.

[63] Daniel M. Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B. Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

A Hyper-parameter selection and pilot results

The β pilot. $\beta = 0.1$ was chosen in October 2026 from $\{0, 0.1, 1\}$ in a pilot on DecisionBench that trained on 645 and tested on 430 of the README's 1,075 test rows, and in a learning-curve study whose test set overlapped the README slice. In that pilot $\beta = 1$ kept about 80% of the gain at 0.8B; in the learning-curve study it kept about a third. The historical grid also ran $\beta = 1$ on the public menu and JevBench at 0.8B with results not included in its original summary: $\beta = 1$ scored 0.584 against 0.590 for $\beta = 0.1$ on the public menu and 0.548 against 0.517 on JevBench (released model 0.550). An independent reproduction of the historical DecisionBench study on another machine recovered the headline gain (+10.4 against +10.2) but not the KL pilot's gap ($\beta = 0$: 0.566 against 0.571 for $\beta = 0.1$, where the pilot had 0.536 against 0.567). The current grid did not retune this choice; $\beta = 0.1$ is used throughout and the no-KL ablation has intervals.

B Reproducing the results

`experiments/peekaboo_remaining/run.py--gridv2` runs the whole grid: it reads every benchmark once through every model and caches the backbone features and the official records, replays each released model against its official prediction call, then runs every method and seed and writes `output_v2/TABLES.md`. Finished steps are skipped, so the run can be resumed. `analysis.py` computes every interval, control and statistic in this paper from the per-row records, and `paper/make_figures.py` produces every table, figure and number macro. The historical grid (`--grid v1`) and its results under `output/` are kept for comparison; its DecisionBench learning-curve study is run by `experiments/vega_policy/reproduce.py`.

How the grid was run. The v2 plan was written for 540 steps. On one A40 the Laya encoder runs took 20 to 45 minutes each and the Laya head runs, which re-run the frozen encoder on every pass, more than an hour on the larger pools, so the plan was trimmed four times while the grid ran (repository commits `50ca000`, `73f3ed5`, `4dbf2e8`, `75a990e`): fewer encoder seeds, the Laya head arms cut to one pass, the realistic-feedback and no-KL arms cut to one seed outside DecisionBench, and no pooled variant. Steps that had already finished were kept, and the committed plan lists exactly the 314 steps that ran. No step was rerun to change a result; GPU out-of-memory failures were retried with the same code and seed.

Checks. Replaying each released model on its cached inputs matches the official prediction call to 0 for Vega and to 5×10^{-5} for Laya (one question per forward pass, as Laya's own call runs it). The released models reproduce their README figures: Vega 0.8B and 4B on the public menu at 0.556 and 0.718 against the README's 0.556 and 0.717, DecisionBench at 0.509 and 0.636 against 0.509 and 0.633; Laya and laya-typed-decisions on typed-decisions at 0.362 and 0.761 against 0.362 and 0.766. Results on the same seed differ across machines by up to 3 points for a single run (floating-point non-determinism in feature extraction and training kernels); these historical reruns are not an independent reproduction of the complete current grid.

C Training seed coverage

Table 12: Independent training runs available in the current grid, counted from per-seed result files. Grouped methods have identical seed sets. A dash means not run. Early stopping reuses the labelled run; it is not another seed. Released models, TTT and historical head-only eight-pass runs are not included.

model	arms	public	JevBench	DecisionB.	typed-dec.	emotion	AG News
Vega 0.8B	feedback and labelled, 8 / 1 passes	3	3	3	3	3	3
Vega 0.8B	logged and damaged feedback	–	–	3	1	1	–
Vega 0.8B	no KL	1	3	3	1	1	1
Vega 0.8B	shipped adapters off	–	–	–	3	–	–
Vega 4B	feedback and labelled, 8 / 1 passes	3	3	3	3	3	3
Vega 4B	logged and damaged feedback	–	–	3	1	1	–
Vega 4B	no KL	1	1	3	1	1	1
Vega 4B	shipped adapters off	–	–	–	3	–	–
Laya	encoder, feedback 8 / 1 and labelled 8	3	3	2	1	1	–
Laya	head, feedback and labelled 1 pass	1	1	1	3	–	–

Table 13: Accuracy on each of the 43 public tasks (mean over seeds). key: the early-stopped answer-key adapter / head; Pk8, Pk1: Peekaboo with eight passes, one pass; f-t: laya-typed-decisions; -enc: encoder unfrozen. * marks the two tasks with a construction artifact. Bold: the best method for that model on that task.

task	Vega 0.8B					Vega 4B					Laya			
	base	TTT	key	Pk8	Pk1	base	TTT	key	Pk8	Pk1	base	f-t	key-enc	Pk-enc
aegis_prompt_safety	0.50	0.50	0.64	**0.64**	0.57	0.60	0.60	**0.79**	0.74	0.73	0.37	0.43	0.68	**0.70**
ag_news	**0.97**	**0.97**	**0.97**	**0.97**	**0.97**	0.95	0.95	**0.97**	0.95	0.97	**0.98**	**0.98**	0.97	0.97
anli_r3	0.28	0.28	**0.39**	0.37	0.32	0.35	**0.52**	0.43	0.40	0.39	**0.43**	0.40	0.42	0.42
arc_challenge	0.48	0.48	0.44	**0.50**	0.48	0.80	0.80	0.80	0.81	**0.83**	0.28	**0.37**	0.34	**0.37**
arc_easy	**0.60**	**0.60**	0.55	0.55	0.58	0.92	0.92	**0.95**	0.94	0.93	0.52	0.57	**0.57**	0.57
banking77	0.68	0.68	**0.72**	0.66	0.62	0.73	0.73	**0.74**	0.69	0.69	0.42	0.40	**0.48**	0.37
boolq	0.75	0.75	0.76	**0.82**	0.79	0.82	0.82	0.84	0.84	**0.91**	**0.87**	0.85	0.78	0.83
casehold	0.35	0.35	**0.49**	0.46	0.42	0.55	0.55	**0.63**	0.61	0.63	0.38	0.38	0.38	**0.45**
climate_fever	0.35	0.38	**0.47**	0.38	0.36	0.45	0.38	**0.48**	0.46	**0.48**	0.33	**0.35**	0.25	0.30
clinc150_oos	**0.68**	**0.68**	0.62	0.62	0.61	**0.92**	**0.92**	0.83	0.78	0.83	0.00	**0.32**	0.00	0.00
commonsenseqa	0.35	0.35	**0.41**	0.38	0.35	0.53	0.53	0.59	0.59	**0.60**	0.45	0.45	**0.47**	0.45
cruxeval_output	0.65	0.65	**0.87**	0.79	0.82	0.85	0.92	**0.94**	0.92	**0.94**	0.78	0.73	0.79	**0.81**
dbpedia14	0.60	**0.85**	0.77	0.67	0.72	0.90	**0.97**	0.92	0.93	0.90	0.85	0.83	0.91	**0.93**
emotion	**0.55**	**0.55**	0.43	0.53	0.54	**0.63**	0.45	0.52	0.52	0.61	**0.57**	**0.57**	0.53	0.53
enron_spam	**1.00**	**1.00**	**1.00**	0.99	**1.00**	**1.00**	**1.00**	**1.00**	**1.00**	0.99	0.92	0.90	**0.93**	0.92
glue_rte	**0.62**	**0.62**	0.61	**0.62**	0.58	0.87	0.85	0.87	0.86	**0.89**	**0.83**	0.82	0.82	0.80
hellaswag	0.25	0.25	0.27	**0.32**	0.28	**0.65**	**0.65**	0.58	0.53	0.59	0.37	0.37	0.39	**0.39**
jailbreak_cls	**0.98**	**0.98**	**0.98**	0.97	0.96	**1.00**	**1.00**	0.99	0.98	0.99	0.98	0.98	0.98	**0.99**
ledgar	0.33	0.33	**0.59**	0.40	0.39	0.63	0.63	0.63	0.63	**0.66**	0.25	0.18	**0.47**	0.33
massive_intent_en	0.38	0.38	**0.53**	0.43	0.43	0.60	0.60	**0.64**	0.59	0.60	0.38	0.42	**0.47**	0.38
medmcqa	0.18	0.18	0.19	0.28	**0.30**	0.48	0.48	**0.49**	**0.49**	0.48	**0.35**	**0.35**	0.32	0.31
medqa_usmle	0.31	0.31	0.27	0.29	**0.31**	0.51	0.51	**0.55**	0.53	**0.55**	0.20	**0.29**	0.23	0.28
mmlu	**0.37**	**0.37**	0.32	0.32	0.31	**0.73**	**0.73**	0.63	0.57	0.61	0.25	**0.35**	0.31	0.32
mmlu_pro	0.20	0.13	0.16	0.21	**0.24**	**0.40**	**0.40**	0.39	**0.40**	0.38	**0.18**	**0.18**	**0.18**	0.18
openbookqa	**0.58**	**0.58**	0.53	0.54	0.58	**0.83**	**0.83**	0.79	0.78	0.79	0.40	0.48	**0.53**	0.49
prompt_injections	0.75	**0.96**	0.87	0.85	0.76	0.85	**0.96**	0.82	0.89	0.90	0.65	0.60	0.81	**0.84**
pubmedqa*	0.13	**0.58**	0.54	0.55	0.54	0.40	**0.62**	0.60	0.56	0.55	0.53	0.55	**0.64**	0.61
race_middle	0.60	0.60	0.54	**0.61**	0.57	0.82	**0.90**	0.82	0.81	0.82	0.58	0.52	**0.59**	0.53
sciq	0.73	0.73	0.67	0.72	**0.76**	0.92	0.92	**0.95**	0.94	0.93	0.52	0.55	**0.56**	0.52
superglue_cb	0.67	0.69	0.68	**0.75**	0.70	0.82	**0.87**	0.81	0.78	0.82	**0.87**	0.84	0.83	0.79
toxic_chat_jailbreak	**0.98**	**0.98**	0.96	0.91	0.92	0.92	0.92	**0.95**	**0.95**	0.94	0.83	**0.90**	0.84	0.87
toxic_chat_toxicity	**0.95**	0.93	0.93	0.90	0.92	**1.00**	0.98	0.98	0.98	0.98	**0.97**	**0.97**	0.94	0.92
trec_coarse	0.53	0.68	**0.72**	0.61	0.60	0.77	**0.88**	0.83	0.77	0.75	0.57	0.52	**0.89**	0.56
truthfulqa_mc1*	0.34	0.34	**0.77**	0.71	0.53	0.64	0.64	**0.73**	0.72	0.65	0.16	0.24	**0.41**	0.38
tweet_hate	**0.65**	0.63	0.62	0.61	0.53	**0.68**	0.60	0.57	0.61	0.59	**0.73**	0.72	0.64	0.67
tweet_irony	**0.78**	0.60	0.71	0.67	0.70	**0.78**	**0.78**	0.74	0.73	0.71	**0.78**	0.74	0.71	0.73
tweet_offensive	0.80	0.80	0.81	0.80	**0.81**	0.78	0.78	**0.85**	0.78	0.80	0.69	0.71	**0.79**	0.78
tweet_sentiment	0.44	0.44	0.45	**0.50**	0.46	**0.58**	**0.58**	0.49	0.56	0.49	**0.56**	0.49	0.45	0.45
tweet_stance_abortion	0.36	0.45	0.48	**0.58**	0.40	**0.55**	0.49	0.48	0.31	0.45	0.31	0.31	0.38	**0.44**
unfair_tos	**0.74**	0.65	0.56	0.57	0.54	**0.88**	0.70	0.82	0.73	0.71	**0.88**	0.84	0.61	0.61
vitaminc	0.47	0.47	**0.54**	0.47	0.53	0.62	0.62	0.63	0.61	**0.64**	0.43	0.33	**0.66**	0.62
winogrande	0.50	0.50	**0.52**	0.49	0.51	0.58	0.58	0.62	**0.63**	**0.63**	0.53	0.50	0.53	**0.57**
yelp_stars	0.50	0.50	**0.59**	0.49	0.45	**0.58**	**0.58**	0.52	0.47	0.52	0.28	0.37	**0.59**	0.38

Table 14: DecisionBench per task (25 test rows each): accuracy of the released model and gain in points of Peekaboo (8 passes), the answer-key adapter and TTT (0.8.0 rule). *: tasks with a construction artifact.

task	Vega 0.8B				Vega 4B			
	base	Pk	key	TTT	base	Pk	key	TTT
tool-action-impact	0.68	+1	+1	+4	0.60	+7	+15	+12
agent-skill-ranking	0.40	+16	+8	+0	0.64	+1	+0	+0
tool-call-routing	0.52	+0	-5	+0	0.52	+5	+0	+0
tool-safety-gate	0.32	+61	+64	+60	0.96	-3	+4	+0
runtime-outcome-verification	0.48	-3	-3	+4	0.44	+3	+0	+8
browser-action-selection	0.52	+8	+16	+0	0.76	-3	-9	+0
browser-target-selection	0.12	+7	+5	+0	0.32	+3	-5	+0
codebase-result-ranking	0.48	+7	-3	+0	0.52	+4	-7	+0
semantic-line-match	0.96	+0	-5	+0	0.96	-1	-5	-4
agent-readiness-assessment	0.72	+8	+12	+0	0.80	+11	+11	+0
agent-call-retention	0.92	+1	-4	-4	0.68	+25	+32	+28
agent-result-retention	1.00	+0	+0	+0	0.92	+8	+8	+0
database-row-classification	0.44	+7	-5	+0	0.40	+4	-3	+0
relevance_score	0.36	-7	+5	+8	0.52	-8	-8	-4
claim-evidence-verification	0.52	+8	+5	-4	0.80	+5	+0	-12
trading-action-selection	0.40	+5	-3	+0	0.36	+9	+9	+0
finqa-numerical-reasoning*	0.16	+39	+49	+0	0.44	+13	+21	+0
route	0.40	+4	+8	+0	0.76	-3	-7	+0
game-goal-selection	0.40	+17	+11	+0	0.52	+7	+12	+0
platformer-control-selection	0.16	+11	+5	+4	0.36	-4	+8	+4
strategy-command-selection	0.16	+17	+19	+0	0.32	+3	+8	+0
route	0.76	-9	+0	+0	0.80	+1	+0	+0
home-alert-triage	0.44	+37	+40	+0	0.64	+24	+23	+0
graph-edge-selection	0.32	+19	+20	+0	0.80	+7	-4	+0
workflow_decision	0.48	+11	+8	+20	0.84	+4	-1	+8
workflow_decision	0.72	+3	+3	+4	0.88	-4	-4	-12
folio-logical-inference	0.44	-9	-9	-24	0.60	-5	-9	-24
musique-evidence-sufficiency	0.60	-8	-1	-24	0.60	-3	-7	-24
musique-multihop-answer	0.52	+27	+33	+0	0.40	+27	+33	+0
contains_threat*	0.76	+24	+24	+0	1.00	+0	+0	+0
toxicity_severity*	0.20	+49	+48	+48	0.20	+51	+49	+24
bounded_value	0.88	+1	+4	+0	1.00	+0	-3	+0
canonical_entity	0.72	+21	+21	+0	0.96	-4	-5	+0
drone-tactical-action	0.16	+36	+52	+0	0.64	+3	+11	+0
robot-skill-selection	0.40	+7	+13	+0	0.48	+1	-7	+0
search-plan-routing	0.56	+19	+15	+0	0.72	+11	+4	+0
social-post-filtering	0.56	+16	+16	+12	0.84	+4	+1	-12
tool_route	0.84	+4	+1	+0	0.92	+0	+0	+0
code-change-risk	0.44	+16	+21	+0	0.52	+8	+12	-4
tax-document-page-classification	0.04	+1	+12	+0	0.12	+0	+3	+0
patent_section	0.40	-4	-1	-12	0.28	+8	+16	+4
relevance	0.76	-4	+4	-12	0.76	+4	+0	+8
dom-ad-detection	0.76	+21	+24	+8	0.76	+24	+24	+16

D Per-task results

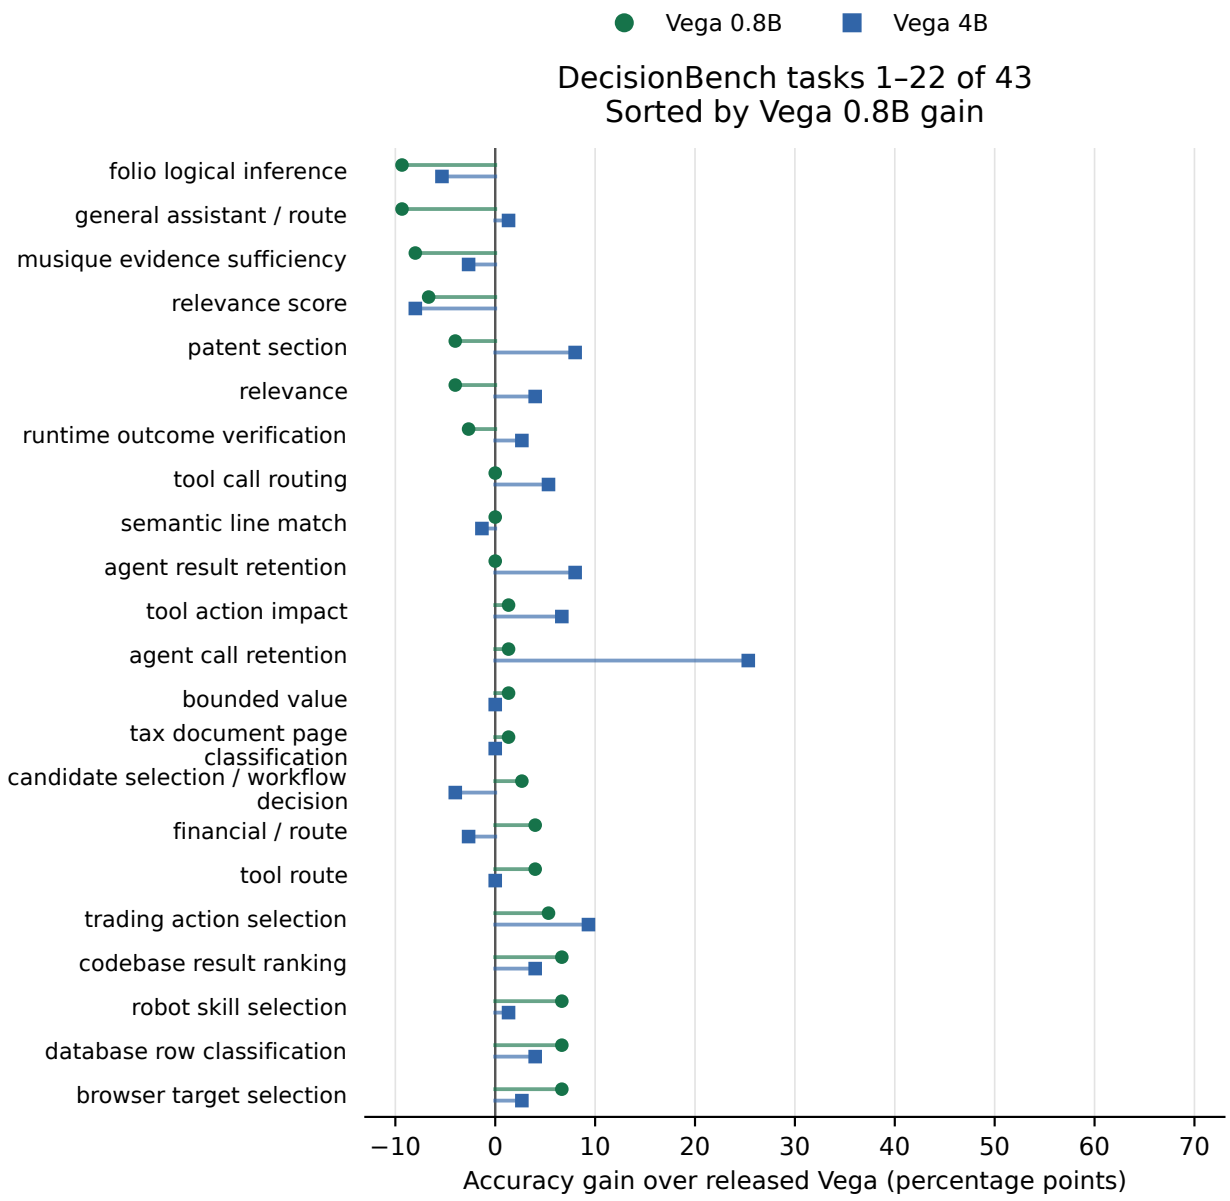


Figure 11: DecisionBench task gains, tasks 1–22 of 43, sorted by the Vega 0.8B gain. Horizontal labels identify each task; circles are Vega 0.8B and squares Vega 4B. Endpoints show mean accuracy gains over the released model; lines connect each point to zero and are not confidence intervals. Each task has 25 test rows. The second panel is Figure 12; both use the same horizontal scale. * marks identified benchmark shortcuts.

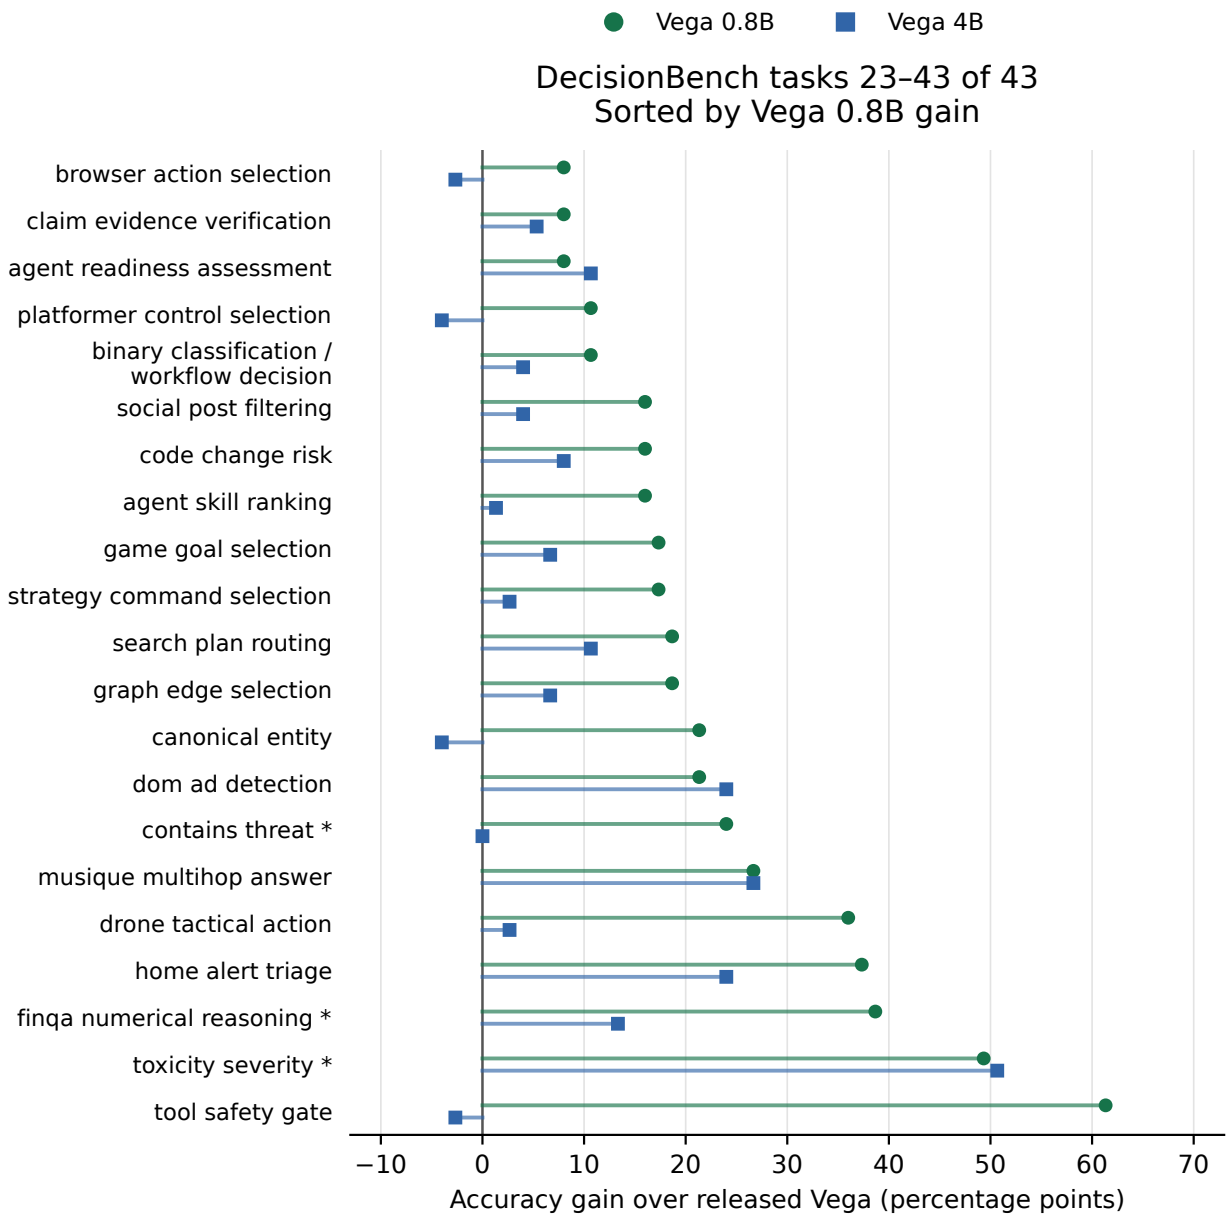


Figure 12: DecisionBench task gains continued, tasks 23–43 of 43. Order, symbols and horizontal scale match Figure 11. * marks identified benchmark shortcuts.

E Hyper-parameters

add-on, Vega	gated LoRA, rank 8, on seven engine layers (`vegaml.attach_adapter`); 409k / 857k parameters
add-on, Laya	trainable copy of type embedding, 2-layer decision head and scorer (26.5M); encoder frozen, or a trained copy ("encoder unfrozen", 421M in all)
optimiser	Adam; lr 10^{-3} (Vega adapter), 10^{-4} (Laya head), 2.5×10^{-5} (Laya encoder); grad-norm clip 1.0
schedule	8 passes (or 1), batches of 32 (encoder: micro-batches of 8, one update per 32 rows)
reward	$r = \mathbf{1}[\text{sampled answer} = \text{gold}]$; baseline $b \leftarrow 0.9\, b + 0.1\, \bar{r}$, $b_0 = 0.5$
KL term	$\beta = 0.1$ for every trained arm (ablation: 0)
answer-key arms	cross-entropy on the gold answer, same KL term; early stopping picks the epoch with the best argmax accuracy on the calibration rows
logged arms	best answer: $-r \log \pi(a) - (1 - r) \log(1 - \pi(a))$ on the released model's argmax, rated once; sampled: clipped importance weights (cap 10) on one answer sampled from the released model
honesty fix	T on a 61-point log grid in $[0.25, 32]$, maximising the Bernoulli likelihood of verdicts on sampled answers from the calibration rows; label arms: NLL temperature scaling on the same rows
switch	keep the add-on if its best answers are right more often than the released model's on the calibration rows
label prior	one log-offset per (task, option text); same loop, lr from $\{10^{-3}, 10^{-2}, 3 \times 10^{-2}, 10^{-1}\}$ chosen on the calibration rows' verdicts
splits	feedback rows 80 / 20 train / calibration per task by a fixed permutation; JevBench 5-fold by paraphrase group, calibration fifth by paraphrase group
intervals	2,000 bootstrap replicates resampling seeds and row clusters; Holm within families; non-inferiority margin 2 points
seeds	0, 1, 2